

# *Dynamic Memory Allocation (2)*

Karthik Dantu


Ethan Blanton

Computer Science and Engineering


University at Buffalo

`kdantu@buffalo.edu`


Slides adapted from CMU 15-213: CSAPP course



# Review: Dynamic Memory Allocation




- Programmers use *dynamic memory allocators* (such as `malloc`) to acquire virtual memory (VM) at run time.
  - for data structures whose size is only known at runtime
- Dynamic memory allocators manage an area of process VM known as the *heap*.




# Review: Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks




Need to tag each block as allocated/free

- Method 2: *Explicit list* among the free blocks using pointers



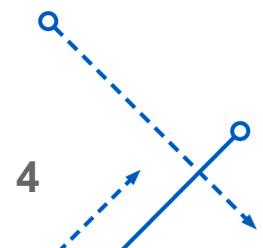
Need space for pointers

- Method 3: *Segregated free list*
  - Different free lists for different size classes
- Method 4: *Blocks sorted by size*
  - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key



# Review: Implicit Lists Summary


- Implementation: very simple
- Allocate cost:
  - linear time worst case
- Free cost:
  - constant time worst case
  - even with coalescing
- Memory Overhead:
  - Depends on placement policy
  - Strategies include first fit, next fit, and best fit
- Not used in practice for malloc/free because of linear-time allocation
  - used in many special purpose applications
- However, the concepts of splitting and boundary tag coalescing are general to *all* allocators



# Today


---

- Explicit free lists
- Segregated free lists
- Garbage collection
- Memory-related perils and pitfalls




# Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks




- Method 2: *Explicit list* among the free blocks using pointers



- Method 3: *Segregated free list*
  - Different free lists for different size classes
- Method 4: *Blocks sorted by size*
  - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

# Explicit Free Lists




- Maintain list(s) of *free* blocks, not *all* blocks
  - Luckily we track only free blocks, so we can use payload area
  - The “next” free block could be anywhere
    - So we need to store forward/back pointers, not just sizes
  - Still need boundary tags for coalescing
    - To find adjacent blocks according to memory order

# Explicit Free Lists

- Logically:




- Physically: blocks can be in any order




# Allocating From Explicit Free Lists

conceptual graphic

*Before*



*After*




*(with splitting)*

= `malloc(...)`

# Freeing With Explicit Free Lists

- **Insertion policy:** Where in the free list do you put a newly freed block?
- **Unordered**
  - LIFO (last-in-first-out) policy
    - Insert freed block at the beginning of the free list
  - FIFO (first-in-first-out) policy
    - Insert freed block at the end of the free list
  - **Pro:** simple and constant time
  - **Con:** studies suggest fragmentation is worse than address ordered
- **Address-ordered policy**
  - Insert freed blocks so that free list blocks are always in address order:  
 $addr(prev) < addr(curr) < addr(next)$
  - **Con:** requires search
  - **Pro:** studies suggest fragmentation is lower than LIFO/FIFO




# Freeing With a LIFO Policy (Case 1)




conceptual graphic

*Before*



- Insert the freed block at the root of the list


*After*




# Freeing With a LIFO Policy (Case 2)



conceptual graphic

*Before*


- Splice out adjacent successor block, coalesce both memory blocks, and insert the new block at the root of the list

*After*


# Freeing With a LIFO Policy (Case 3)




conceptual graphic

*Before*


- Splice out adjacent predecessor block, coalesce both memory blocks, and insert the new block at the root of the list

*After*

# Freeing With a LIFO Policy (Case 4)



- Splice out adjacent predecessor and successor blocks, coalesce all 3 blocks, and insert the new block at the root of ~~the list~~




# Explicit List Summary

- Comparison to implicit list:
  - Allocate is linear time in number of **free** blocks instead of **all** blocks
    - **Much faster** when most of the memory is full
  - Slightly more complicated allocate and free because need to splice blocks in and out of the list
  - Some extra space for the links (2 extra words needed for each block)
    - Does this increase internal fragmentation?

# Today


---

- Explicit free lists
- Segregated free lists
- Garbage collection
- Memory-related perils and pitfalls



# Segregated List (Seglist) Allocators


- Each *size class* of blocks has its own free list



- Often have separate classes for each small size
- For larger sizes: One class for each size  $[2^i + 1, 2^{i+1}]$

# Seglist Allocator

- Given an array of free lists, each one for some size class
- To allocate a block of size  $n$ :
  - Search appropriate free list for block of size  $m > n$  (i.e., first fit)
  - If an appropriate block is found:
    - Split block and place fragment on appropriate list
    - If no block is found, try next larger class
  - Repeat until block is found
- If no block is found:
  - Request additional heap memory from OS (using **sbrk()**)
  - Allocate block of  $n$  bytes from this new memory
  - Place remainder as a single free block in appropriate size class.



# Seglist Allocator (cont.)

- To free a block:
  - Coalesce and place on appropriate list
- Advantages of seglist allocators vs. non-seglist allocators (both with first-fit)
  - Higher throughput
    - log time for power-of-two size classes vs. linear time
  - Better memory utilization
    - First-fit search of segregated free list approximates a best-fit search of entire heap.
    - Extreme case: Giving each block its own size class is equivalent to best-fit.

# Today

---


- Explicit free lists
- Segregated free lists
- Garbage collection
- Memory-related perils and pitfalls

# Implicit Memory Management: Garbage Collection

- *Garbage collection*: automatic reclamation of heap-allocated storage—application never has to explicitly free memory


```
void foo() {  
    int *p = malloc(128);  
    return; /* p block is now garbage */  
}
```

- Common in many dynamic languages:
  - Python, Ruby, Java, Perl, ML, Lisp, Mathematica
- Variants (“conservative” garbage collectors) exist for C and C++
  - However, cannot necessarily collect all garbage



# Garbage Collection

- How does the memory manager know when memory can be freed?
  - In general we cannot know what is going to be used in the future since it depends on conditionals
  - But we can tell that certain blocks cannot be used if there are no pointers to them
- Must make certain assumptions about pointers
  - Memory manager can distinguish pointers from non-pointers
  - All pointers point to the start of a block
  - Cannot hide pointers
    - (e.g., by coercing them to an `int`, and then back again)




# Classical GC Algorithms

- Mark-and-sweep collection (McCarthy, 1960)
  - Does not move blocks (unless you also “compact”)
- Reference counting (Collins, 1960)
  - Does not move blocks (not discussed)
- Copying collection (Minsky, 1963)
  - Moves blocks (not discussed)
- Generational Collectors (Lieberman and Hewitt, 1983)
  - Collection based on lifetimes
    - Most allocations become garbage very soon
    - So focus reclamation work on zones of memory recently allocated
- For more information:  
*Jones and Lin, “Garbage Collection: Algorithms for Automatic Dynamic Memory”, John Wiley & Sons, 1996.*



# Memory as a Graph

- We view memory as a directed graph
  - Each block is a node in the graph
  - Each pointer is an edge in the graph
  - Locations not in the heap that contain pointers into the heap are called **root** nodes (e.g. registers, locations on the stack, global variables)




A node (block) is **reachable** if there is a path from any root to that node.

Non-reachable nodes are **garbage** (cannot be needed by the application)

# Mark and Sweep Collecting

- Can build on top of malloc/free package
  - Allocate using `malloc` until you “run out of space”
- When out of space:
  - Use extra **mark bit** in the head of each block
  - **Mark:** Start at roots and set mark bit on each reachable block
  - **Sweep:** Scan all blocks and free blocks that are not marked




# Today

---

- Explicit free lists
- Segregated free lists
- Garbage collection
- Memory-related perils and pitfalls

# Memory-Related Perils and Pitfalls

- Dereferencing bad pointers
- Reading uninitialized memory
- Overwriting memory
- Referencing nonexistent variables
- Freeing blocks multiple times
- Referencing freed blocks
- Failing to free blocks



# Dereferencing Bad Pointers

- The classic `scanf` bug

```
int val;  
  
...  
  
scanf("%d", val);
```

# Reading Uninitialized Memory

- Assuming that heap data is initialized to zero

```
/* return y = Ax */
int *matvec(int **A, int *x) {
    int *y = malloc(N*sizeof(int));
    int i, j;

    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            y[i] += A[i][j]*x[j];
    return y;
}
```


- Can avoid by using calloc

# Overwriting Memory

- Allocating the (possibly) wrong sized object

```
int **p;  
  
p = malloc(N*sizeof(int));  
  
for (i=0; i<N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

- Can you spot the bug?



# Overwriting Memory

- Off-by-one errors

```
char **p;  
  
p = malloc(N*sizeof(int *));  
  
for (i=0; i<=N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

```
char *p;  
  
p = malloc(strlen(s));  
strcpy(p, s);
```


# Overwriting Memory

- Not checking the max string size

```
char s[8];
int i;

gets(s); /* reads "123456789" from stdin */
```


- Basis for classic buffer overflow attacks



# Overwriting Memory

- Misunderstanding pointer arithmetic

```
int *search(int *p, int val) {  
  
    while (p && *p != val)  
        p += sizeof(int);  
  
    return p;  
}
```



# Overwriting Memory


- Referencing a pointer instead of the object it points to

```
int *BinheapDelete(int **binheap, int *size) {  
    int *packet;  
    packet = binheap[0];  
    binheap[0] = binheap[*size - 1];  
    *size--;  
    Heapify(binheap, *size, 0);  
    return (packet);  
}
```

- What gets decremented?
  - (See next slide)



# C operators



- $\rightarrow$ ,  $()$ , and  $[]$  have high precedence, with  $*$  and  $\&$  just below
- Unary  $+$ ,  $-$ , and  $*$  have higher precedence than binary forms

# Overwriting Memory

- Referencing a pointer instead of the object it points to

```
int *BinheapDelete(int **binheap, int *size) {  
    int *packet;  
    packet = binheap[0];  
    binheap[0] = binheap[*size - 1];  
    *size--;  
    Heapify(binheap, *size, 0);  
    return (packet);  
}
```

**Operators**

|                                   |     |                 |
|-----------------------------------|-----|-----------------|
| ( ) [ ] -> . ++ --                | *   | & (type) sizeof |
| ! ~ ++ -- + -                     | /*  |                 |
| *                                 | / % |                 |
| +                                 | -   |                 |
| <<                                | >>  |                 |
| < <= > >=                         |     |                 |
| == !=                             |     |                 |
| &                                 |     |                 |
| ^                                 |     |                 |
|                                   |     |                 |
| &&                                |     |                 |
|                                   |     |                 |
| ?:                                |     |                 |
| = += -= *= /= %= &= ^= != <<= >>= |     |                 |
| ,                                 |     |                 |

**Associativity**

|               |
|---------------|
| left to right |
| right to left |
| left to right |
| right to left |
| right to left |
| left to right |

- Same effect as
  - size--;**
- Rewrite as
  - (\*size)--;**

# Referencing Nonexistent Variables

- Forgetting that local variables disappear when a function returns

```
int *foo () {  
    int val;  
  
    return &val;  
}
```

# Freeing Blocks Multiple Times

- Nasty!

```
x = malloc(N*sizeof(int));  
    <manipulate x>  
free(x);  
  
y = malloc(M*sizeof(int));  
    <manipulate y>  
free(x);
```

# Referencing Freed Blocks

- Evil!

```
x = malloc(N*sizeof(int));  
    <manipulate x>  
free(x);  
    ...  
y = malloc(M*sizeof(int));  
for (i=0; i<M; i++)  
    y[i] = x[i]++;
```

# Failing to Free Blocks (Memory Leaks)


- Slow, long-term killer!

```
foo() {  
    int *x = malloc(N*sizeof(int));  
    ...  
    return;  
}
```

# Failing to Free Blocks (Memory Leaks)

- Freeing only part of a data structure

```
struct list {  
    int val;  
    struct list *next;  
};  
  
foo() {  
    struct list *head = malloc(sizeof(struct list));  
    head->val = 0;  
    head->next = NULL;  
    <create and manipulate the rest of the list>  
    ...  
    free(head);  
    return;  
}
```



# Dealing With Memory Bugs

- Debugger: gdb
  - Good for finding bad pointer dereferences
  - Hard to detect the other memory bugs
- Data structure consistency checker
  - Runs silently, prints message only on error
  - Use as a probe to zero in on error
- Binary translator: valgrind
  - Powerful debugging and analysis technique
  - Rewrites text section of executable object file
  - Checks each individual reference at runtime
    - Bad pointers, overwrites, refs outside of allocated block
- glibc malloc contains checking code
  - `setenv MALLOC_CHECK_ 3`

