

Floating Point Numbers

Karthik Dantu

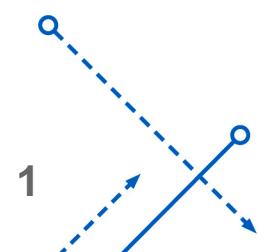
Ethan Blanton

Computer Science and Engineering

University at Buffalo

kdantu@buffalo.edu

Karthik Dantu




Administrivia

- Midterm: Oct 9th (Wednesday) in class
- PA2 out now
- gdb lab

Today: Floating Point


- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary



Fractional binary numbers

- What is 1011.101_2 ?

Fractional Binary Numbers




- Representation

- Bits to right of “binary point” represent fractional powers of 2
- Represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$

Karthik Dantu



Fractional Binary Numbers: Examples

■ Value

$$5 \frac{3}{4} = 23/4$$

$$2 \frac{7}{8} = 23/8$$

$$1 \frac{7}{16} = 23/16$$

Representation

$$101.11_2$$

$$10.111_2$$

$$1.0111_2$$

$$= 4 + 1 + 1/2 + 1/4$$

$$= 2 + 1/2 + 1/4 + 1/8$$

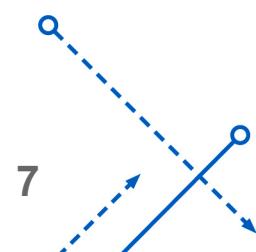
$$= 1 + 1/4 + 1/8 + 1/16$$

■ Observations

- Divide by 2 by shifting right (unsigned)
- Multiply by 2 by shifting left
- Numbers of form $0.111111\dots_2$ are just below 1.0
 - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
 - Use notation $1.0 - \epsilon$


Representable Numbers

- Limitation #1
 - Can only exactly represent numbers of the form $x/2^k$
 - Other rational numbers have repeating bit representations
 - Value Representation
 - $1/3$ $0.0101010101[01]..._2$
 - $1/5$ $0.001100110011[0011]..._2$
 - $1/10$ $0.0001100110011[0011]..._2$
- Limitation #2
 - Just one setting of binary point within the w bits
 - Limited range of numbers (very small values? very large?)



Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary



IEEE Floating Point

- IEEE Standard 754
 - Established in 1985 as uniform standard for floating point arithmetic
 - Before that, many idiosyncratic formats
 - Supported by all major CPUs
 - Some CPUs don't implement IEEE 754 in full
 - e.g., early GPUs, Cell BE processor
- Driven by numerical concerns
 - Nice standards for rounding, overflow, underflow
 - Hard to make fast in hardware
 - **Numerical analysts** predominated over **hardware designers** in defining standard

Floating Point Representation

- Numerical Form:


$$(-1)^s \ M \ 2^E$$

- Sign bit **s** determines whether number is negative or positive
- Significand **M** normally a fractional value in range [1.0,2.0).
- Exponent **E** weights value by power of two
- Encoding
 - MSB **s** is sign bit **s**
 - exp field encodes **E** (but is not equal to E)
 - frac field encodes **M** (but is not equal to M)



Example:

$$15213_{10} = (-1)^0 \times 1.110110110110_2 \times 2^{13}$$



Precision options


- Single precision: 32 bits
 ≈ 7 decimal digits, $10^{\pm 38}$




- Double precision: 64 bits
 ≈ 16 decimal digits, $10^{\pm 308}$



- Other formats: half precision, quad precision




Three “kinds” of floating point numbers



“Normalized” Values

$$v = (-1)^s M \cdot 2^E$$

- When: $\exp \neq 000\dots0$ and $\exp \neq 111\dots1$
- Exponent coded as a biased value: $E = \exp - \text{Bias}$
 - \exp : unsigned value of exp field
 - Bias = $2^{k-1} - 1$, where k is number of exponent bits
 - Single precision: 127 (\exp : 1...254, E: -126...127)
 - Double precision: 1023 (\exp : 1...2046, E: -1022...1023)
- Significand coded with implied leading 1: $M = 1.\text{xxx...x}_2$
 - xxx...x : bits of frac field
 - Minimum when $\text{frac}=000\dots0$ ($M = 1.0$)
 - Maximum when $\text{frac}=111\dots1$ ($M = 2.0 - \epsilon$)
 - Get extra leading bit for “free”



Normalized Encoding Example

$$v = (-1)^s M \cdot 2^E$$
$$E = \text{exp} - \text{Bias}$$

- Value: float F = 15213.0;
 - $15213_{10} = 11101101101101_2$
 $= 1.1101101101101_2 \times 2^{13}$

- Significand

$$M = 1.\underline{1101101101101}_2$$

$$\text{frac} = \underline{1101101101101}0000000000_2$$

- Exponent

$$E = 13$$

$$Bias = 127$$

$$\text{exp} = 140 = 10001100_2$$

- Result:

0	10001100	110110110110100000000000
s	exp	frac


Denormalized Values

$$v = (-1)^s M \cdot 2^E$$
$$E = 1 - \text{Bias}$$

- Condition: $\exp = 000\dots0$
- Exponent value: $E = 1 - \text{Bias}$ (instead of $\exp - \text{Bias}$) (why?)
- Significand coded with implied leading 0: $M = 0.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
- Cases
 - $\exp = 000\dots0, \text{frac} = 000\dots0$
 - Represents zero value
 - Note distinct values: +0 and -0 (why?)
 - $\exp = 000\dots0, \text{frac} \neq 000\dots0$
 - Numbers closest to 0.0
 - Equispaced

Special Values

- Condition: `exp = 111...1`
- Case: `exp = 111...1, frac = 000...0`
 - **Represents value ∞ (infinity)**
 - Operation that overflows
 - Both positive and negative
 - E.g., $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$
- Case: `exp = 111...1, frac \neq 000...0`
 - **Not-a-Number (NaN)**
 - Represents case when no numeric value can be determined
 - E.g., $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$



C float Decoding Example

$$v = (-1)^s M 2^E$$

$$E = \text{exp} - \text{Bias}$$

$$\text{Bias} = 2^{k-1} - 1 = 127$$

float: 0xC0A00000

binary:



1

8-bits

23-bits

E =

S =

M =

$v = (-1)^s M 2^E =$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

C float Decoding Example #1

$$v = (-1)^s M 2^E$$

$$E = \text{exp} - \text{Bias}$$

float: 0xC0A00000

binary: 1100 0000 1010 0000 0000 0000 0000 0000



$$E =$$

$$S =$$

$$M = 1.$$

$$v = (-1)^s M 2^E =$$

Hex Decimal Binary

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

C float Decoding Example #1


$$v = (-1)^s M 2^E$$

$$E = \text{exp} - \text{Bias}$$

$$\text{Bias} = 2^{k-1} - 1 = 127$$

float: 0xC0A00000

binary: 1100 0000 1010 0000 0000 0000 0000 0000



1

8-bits

23-bits

$$E = \text{exp} - \text{Bias} = 129 - 127 = 2 \text{ (decimal)}$$

S = 1 \rightarrow negative number

$$\begin{aligned} M &= 1.010 0000 0000 0000 0000 0000 \\ &= 1 + 1/4 = 1.25 \end{aligned}$$

$$v = (-1)^s M 2^E = (-1)^1 * 1.25 * 2^2 = -5$$

Hex Decimal Binary

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

C float Decoding Example #2

$$v = (-1)^s M 2^E$$
$$E = 1 - \text{Bias}$$

float: 0x001c0000

binary: 0000 0000 0001 1100 0000 0000 0000 0000



1

8-bits

23-bits

$$E =$$

$$S =$$

$$M = 0.$$

$$v = (-1)^s M 2^E =$$

Hex Decimal Binary

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

C float Decoding Example #2


$$v = (-1)^s M 2^E$$

$$E = 1 - \text{Bias}$$

$$\text{Bias} = 2^{k-1} - 1 = 127$$

float: 0x001C0000

binary: 0000 0000 0001 1100 0000 0000 0000 0000



$$E = 1 - \text{Bias} = 1 - 127 = -126 \text{ (decimal)}$$

S = 0 -> positive number

$$M = 0.001 1100 0000 0000 0000 0000$$

$$= 1/8 + 1/16 + 1/32 = 7/32 = 7 \cdot 2^{-5}$$


$$v = (-1)^s M 2^E = (-1)^0 * 7 \cdot 2^{-5} * 2^{-126} = 7 \cdot 2^{-131}$$

$$\approx 2.571393892 \times 10^{-39}$$

Hex Decimal Binary

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111


Visualization: Floating Point Encodings



Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary

Tiny Floating Point Example



- 8-bit Floating Point Representation
 - the sign bit is in the most significant bit
 - the next four bits are the **exp**, with a bias of 7
 - the last three bits are the **frac**
- Same general form as IEEE Format
 - normalized, denormalized
 - representation of 0, NaN, infinity

Dynamic Range (s=0 only)

	s	exp	frac	E	Value	
Denormalized numbers	0	0000	000	-6	0	
	0	0000	001	-6	$1/8 * 1/64 = 1/512$	closest to zero
	0	0000	010	-6	$2/8 * 1/64 = 2/512$	$(-1)^0 (0+1/4) * 2^{-6}$
	...					
	0	0000	110	-6	$6/8 * 1/64 = 6/512$	
	0	0000	111	-6	$7/8 * 1/64 = 7/512$	largest denorm
	0	0001	000	-6	$8/8 * 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 * 1/64 = 9/512$	$(-1)^0 (1+1/8) * 2^{-6}$
	...					
	0	0110	110	-1	$14/8 * 1/2 = 14/16$	
Normalized numbers	0	0110	111	-1	$15/8 * 1/2 = 15/16$	closest to 1 below
	0	0111	000	0	$8/8 * 1 = 1$	
	0	0111	001	0	$9/8 * 1 = 9/8$	closest to 1 above
	0	0111	010	0	$10/8 * 1 = 10/8$	
	...					
	0	1110	110	7	$14/8 * 128 = 224$	
	0	1110	111	7	$15/8 * 128 = 240$	largest norm
	0	1111	000	n/a	inf	


$$v = (-1)^s M 2^E$$

norm: $E = \text{exp} - \text{Bias}$


denorm: $E = 1 - \text{Bias}$

Distribution of Values

- 6-bit IEEE-like format
 - $e = 3$ exponent bits
 - $f = 2$ fraction bits
 - Bias is $2^{3-1}-1 = 3$





- Notice how the distribution gets denser toward zero.



Distribution of Values (close-up view)

- 6-bit IEEE-like format
 - $e = 3$ exponent bits
 - $f = 2$ fraction bits
 - Bias is 3



Special Properties of the IEEE Encoding

- FP Zero Same as Integer Zero
 - All bits = 0
- Can (Almost) Use Unsigned Integer Comparison
 - Must first compare sign bits
 - Must consider $-0 = 0$
 - NaNs problematic
 - Will be greater than any other values
 - What should comparison yield? The answer is complicated.
 - Otherwise OK
 - Denorm vs. normalized
 - Normalized vs. infinity



Floating Point Operations: Basic Idea

- $x +_f y = \text{Round}(x + y)$
- $x \times_f y = \text{Round}(x \times y)$
- Basic idea
 - First **compute exact result**
 - Make it fit into desired precision
 - Possibly overflow if exponent too large
 - Possibly **round to fit into `frac`**



Rounding

- Rounding Modes (illustrate with \$ rounding)

	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
• Towards zero	\$1	\$1	\$1	\$2	-\$1
• Round down ($-\infty$)	\$1	\$1	\$1	\$2	-\$2
• Round up ($+\infty$)	\$2	\$2	\$2	\$3	-\$1
• Nearest Even* (default)		\$1	\$2 [↑]	\$2 [↓]	\$2 [↓]
	-\$2				

*Round to nearest, but if half-way in-between then round to nearest even

Closer Look at Round-To-Even

- Default Rounding Mode
 - Hard to get any other kind without dropping into assembly
 - C99 has support for rounding mode management
 - All others are statistically biased
 - Sum of set of positive numbers will consistently be over- or under-estimated
- Applying to Other Decimal Places / Bit Positions
 - When exactly halfway between two possible values
 - Round so that least significant digit is even
 - E.g., round to nearest hundredth

7.8949999	7.89	(Less than half way)
7.8950001	7.90	(Greater than half way)
7.8950000	7.90	(Half way—round up)
7.8850000	7.88	(Half way—round down)



Rounding Binary Numbers

- Binary Fractional Numbers
 - “Even” when least significant bit is 0
 - “Half way” when bits to right of rounding position = $100\dots_2$

- Examples
 - Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded Value
$2 \frac{3}{32}$	$10.00\textcolor{red}{011}_2$	10.00_2	($<1/2$ —down)	2
$2 \frac{3}{16}$	$10.00\textcolor{red}{110}_2$	10.01_2	($>1/2$ —up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\textcolor{red}{100}_2$	11.00_2	($\frac{1}{2}$ —up)	3
$2 \frac{5}{8}$	$10.10\textcolor{red}{100}_2$	10.10_2	($\frac{1}{2}$ —down)	$2 \frac{1}{2}$

Rounding

1.BBGRXXXX

Guard bit: LSB of result

Round bit: 1st bit removed


Sticky bit: OR of remaining bits

- Round up conditions
 - Round = 1, Sticky = 1 $\rightarrow > 0.5$
 - Guard = 1, Round = 1, Sticky = 0 \rightarrow Round to even

Fraction	GRS	Incr?	Rounded
1.0000000	000	N	1.000
1.1010000	100	N	1.101
1.0001000	010	N	1.000
1.0011000	110	Y	1.010
1.0001010	011	Y	1.001
1.1111100	111	Y	10.000

FP Multiplication

- $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$
- Exact Result: $(-1)^s M 2^E$
 - Sign s: $s_1 \wedge s_2$
 - Significand M: $M_1 \times M_2$
 - Exponent E: $E_1 + E_2$
- Fixing
 - If $M \geq 2$, shift M right, increment E
 - If E out of range, overflow
 - Round M to fit `frac` precision
- Implementation
4 bit significand: $1.010 \times 2^2 \times 1.110 \times 2^3 = 10.0011 \times 2^5$
 - Biggest chore is multiplying significands 1.00011×2^6 and 1.001×2^6



Floating Point Addition

$$\bullet (-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$$

- Assume $E_1 > E_2$

Get binary points lined up


$$\bullet \text{Exact Result: } (-1)^s M 2^E$$

- Sign s , significand M :
 - Result of signed align & add

- Exponent E : E_1

- Fixing

- If $M \geq 2$, shift M right, increment E
- if $M < 1$, shift M left k positions, decrement E by k
- Overflow if E out of range
- Round M to fit `frac` precision



$$\begin{aligned}1.010 * 2^2 + 1.110 * 2^3 &= (0.1010 + 1.1100) * 2^3 \\&= 10.0110 * 2^3 = 1.001\textcolor{red}{10} * 2^4 = 1.010 * 2^4\end{aligned}$$

- Compare to those of Abelian Group
 - Closed under addition?
 - But may generate infinity or NaN
 - Commutative?
 - Yes
 - Associative?
 - Overflow and inexactness of rounding
 - $(3.14+1e10)-1e10 = 0$, $3.14+(1e10-1e10) = 3.14$
 - 0 is additive identity?
 - Yes
 - Every element has additive inverse?
 - Yes, except for infinities & NaNs
 - No
- Monotonicity
 - $a \geq b \Rightarrow a+c \geq b+c?$
 - Except for infinities & NaNs
 - Almost

Mathematical Properties of FP Mult

- Compare to Commutative Ring
 - Closed under multiplication?
 - But may generate infinity or NaN
 - Multiplication Commutative?
 - Multiplication is Associative?
 - Possibility of overflow, inexactness of rounding
 - Ex: $(1e20 * 1e20) * 1e-20 = \inf$, $1e20 * (1e20 * 1e-20) = 1e20$
 - 1 is multiplicative identity?
 - Multiplication distributes over addition?
 - Possibility of overflow, inexactness of rounding
 - $1e20 * (1e20 - 1e20) = 0.0$, $1e20 * 1e20 - 1e20 * 1e20 = \text{NaN}$
 - Monotonicity
 - $a \geq b \ \& \ c \geq 0 \Rightarrow a * c \geq b * c?$
 - Except for infinities & NaNs

Yes

Yes

No

Yes

No

Almost



