

---

# CSE 220: Systems Programming

## 1 - *Introduction*

Karthik Dantu

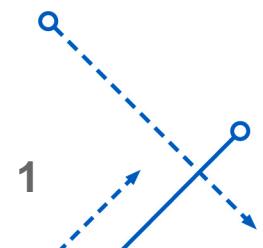
Ethan Blanton

Computer Science and Engineering

University at Buffalo

kdantu@buffalo.edu

Karthik Dantu



# CSE 220: Systems Programming

- **Instructor:** Karthik Dantu (this section)  
Ethan Blanton (another section)
- **Office Hours:** 12:30-2:00 pm (MW)
- **Course Website:** <https://droneslab.github.io/cse220/>
- **Discussion Forum:**  
<https://piazza.com/buffalo/fall2019/cse220/resources>
- **Syllabus:** <https://droneslab.github.io/cse220/lectures/syllabus.pdf>
- **Textbooks:**  
Randal E Bryant and David R O'Hallaron. *Computer Science: A Programmer's Perspective*. Third Edition. Pearson. 2016.  
Brian W Kernighan and Dennis Ritchie. *The C Programming Language*. Second Edition. Prentice Hall. 1988.

# CSE 220: Objectives

- **Objective:** Understand how hardware (processor, memory, GPU, disks, network) and software (OS, compilers, libraries) come together to execute application programs
- **Benefits of CSE 220**

Become better programmers

Identify and eliminate bugs and program bottlenecks efficiently

Understand and tune program performance

Stepping stone for other Systems classes in CS and CE

CSE 421: Operating systems

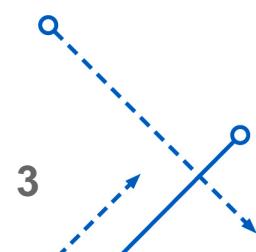
CSE 489: Modern Networking Concepts

CSE 305: Intro to Programming Languages

CSE 341: Computer Organization

CSE 321: Realtime and Embedded Systems

.... and many more



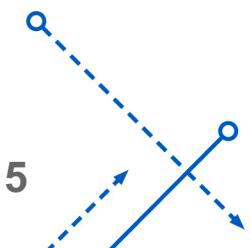
# CSE 220: Content

- Hardware
  - Decades of evolution of processor technology
  - Memory and storage have evolved in parallel
  - I/O (including disk and network) dominate interactivity changing the usage model
- Software
  - OS abstracts the hardware and provides a unified interface (system calls)
  - Libraries provide interfaces for commonly used programming fragments (data structures, math operations, user interaction etc. )
- Application programs
  - Sit on top of all the above hardware/software
  - Execute *higher-level commands*
  - Are typically compiled/interpreted by a compiler/interpreter
- This class helps you understand the interaction of application programs with all the above - and hopefully interest you in delving into looking under the hood of a computer program!

# CSE 220: Expectations

- Probably your toughest CSE class until now - one of the harder classes in general!
- Attendance - Mandatory!
  - I will not re-do lectures
  - Recitations will also not repeat lectures
  - If you skip lectures, expectation is that you'll catch up on your own
- Labs - at least as important as the lectures!
  - Practice what you learn in class
  - Frequent lab exams to test your understanding
  - Significant portion of the grade
- Assistance - ask for help *early* and *often*
- Meet pre-reqs
  - Some programming experience
  - Understand linked lists and object references

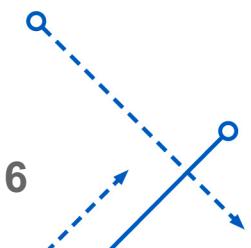
This class elevates expectations from you in terms of your off-class learning. We strongly believe this is the systems-way of thinking – and 220 is designed to inculcate this in you!



# CSE 220: Etiquette

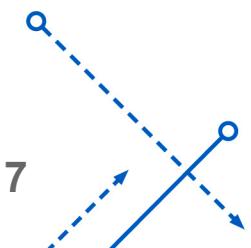
- Attend every class and lab !
- Be respectful to instructors, TAs and classmates
- **Adhere strictly** to the academic integrity policy (more in a bit)

Behave as adults and strive to maximize your and your classmates' learning experience in this course.



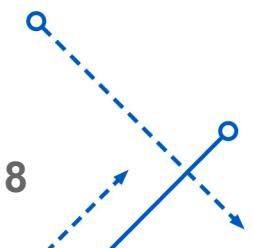
# CSE 220: Ways To Fail

- Missing classes and labs
- Start assignments at the last minute
- Not visiting office hrs
- Not asking questions on Piazza
- Waiting until the deadline to submit for the first time
- Cheat!



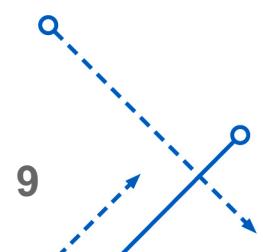
## Academic Integrity

Cheating is submitting any work that you did not perform by yourself as if you did.



# CSE220: Academic Integrity Guidelines

- References (when cited)  
Wikipedia, Wikibooks (or similar): **OK**
- Public Code  
StackExchange/StackOverflow: **Not OK**
- Discussing concepts/ideas with classmates  
“A hash index has O(1) lookups”: **OK** (except exams)
- Sharing code or answers with anyone  
“Just look at how I implemented it”: **Not OK**  
For-hire code: **NOT OK**



# CSE 220: Academic Integrity

- We use sophisticated code checkers such as `moss`
- Trust me – it is better than any disguise you can do in short order!
- We also compare with submissions from previous years, as well as publicly available repos

## Moss Results

Sat Mar 2 20:19:46 PST 2013

Options -l java -d -m 10

CSE 562 Project 1

[ [How to Read the Results](#) | [Tips](#) | [FAQ](#) | [Contact](#) | [Submission Scripts](#) | [Credits](#) ]

File 1	File 2	Lines Matched
(52%)	(52%)	2142
(84%)	84%	1556
(40%)	(40%)	1194
(29%)	(18%)	1163
(13%)	(19%)	822
10%)	(9%)	569
(11%)	(10%)	660
(11%)	(16%)	616
(10%)	(7%)	513
(10%)	8%)	613

## Submission Overlap (Ignoring Library Code)

## Identical Code Structure

The diagram illustrates the identical code structure between two sections of Java code. On the left, a `case` statement handles various operators. On the right, an `operator class` provides the implementation for these operators. Colored arrows and boxes highlight the correspondence between the two sections.

**Code in Case Statement**

```
    }  
    break;  
  
    //Implementation for LEAF operators  
    case LEAF: {  
        switch (q.type) {  
            //Implementation for NULLSOURCE operator  
            case NULLSOURCE: {  
                Datum[] datum = new Datum[0];  
                query_res.data_list.add(datum);  
                break;  
            }  
            //Implementation for SCAN operator  
            case SCAN: {  
  
                ScanNode node = (ScanNode)q;  
                File file = tables.get(node.table).getFile();  
                FileReader fir;  
                BufferedReader reader;  
                Datum[] datum;  
                try {  
  
                    fir = new FileReader(file);  
                    reader = new BufferedReader(fir);  
                    String record;  
                    StringTokenizer st = null;  
                    while ((record = reader.readLine()) != null) {  
                        datum = new Datum[node.schema.size()];  
                        int i = 0;  
                        st = new StringTokenizer(record, ",");  
                        while (st.hasMoreElements()) {  
                            String token = st.nextToken().trim();  
                            switch (node.schema.get(i).type) {  
                                case INT:  
                                    datum[i] = new Datum.Int(Integer.parseInt(t  
                                case FLOAT:  
                                    datum[i] = new Datum.Flt(Float.parseFloat(t  
                                case STRING:  
                                    datum[i] = new Datum.Str(token);  
                                case BOOL:  
                                    datum[i] = new Datum.Bool(new Boolean(token  
                                default:  
                                    break;  
                            }  
                            break;  
                        }  
                    }  
                } catch (Exception e) {  
                    System.out.println("Error reading file " + file.getName());  
                }  
            }  
        }  
    }  
}
```

**Code in “Operator Class”**

```
import edu.buffalo.cse.sql.Schema.TableFromFile;  
import edu.buffalo.cse.sql.data.Datum;  
import edu.buffalo.cse.sql.plan.PlanNode;  
import edu.buffalo.cse.sql.plan.ScanNode;  
  
public class ScanMethod extends UnionMethod {  
  
    public static void scanop(Map<String, Schema.TableFromFile> tables, PlanNode plan  
        throws NumberFormatException, IOException {  
  
        ScanNode leafnode = (ScanNode) q;  
        File file = tables.get(leafnode.table).getFile();  
        FileReader fir;  
        BufferedReader reader;  
        Datum[] data;  
        int sizeofschemas= leafnode.schema.size();  
  
        fir = new FileReader(file);  
        reader = new BufferedReader(fir);  
        String records;  
        StringTokenizer st = null;  
        while ((records = reader.readLine()) != null) {  
            data = new Datum[leafnode.schema.size()];  
            int i = 0;  
            st = new StringTokenizer(records, ",");  
            while (st.hasMoreElements()) {  
                String token = st.nextToken().trim();  
                switch (leafnode.schema.get(i).type) {  
                    case INT:  
                        data[i] = new Datum.Int(Integer.parseInt(t  
                    case FLOAT:  
                        data[i] = new Datum.Flt(Float.parseFloat(t  
                    case STRING:  
                        data[i] = new Datum.Str(token);  
                    case BOOL:  
                        data[i] = new Datum.Bool(new Boolean(token  
                    default:  
                        break;  
                }  
                i++;  
            }  
        }  
        reparse.schemaVar.add(data);  
    }  
}  
fir.close();
```

Code in Case Statement

Code in “Operator Class”

# CSE 220: Academic Integrity Policy

- First offense

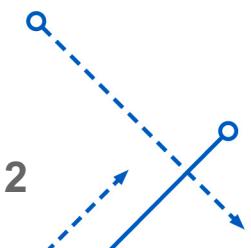
Zero on the assignment, lower grade, or an F in the class (instructor judgement)

Your name reported to university black list that will follow you through your time at UB

If your name is already on black-list, you get an F with note – **you will fail the class and transcript will reflect your cheating**

- Share code, share blame

If someone else submits your code as their own, **you will be penalized as well**

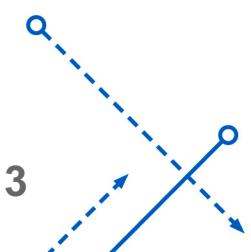


# CSE 220: AI Best Policies

- Be careful with your code – including permissions on shared UB filesystems, GitHub, Bitbucket etc.
- **Don't look** at someone else's code!
- Cite liberally
- Check with department[1] and university[2] policies
- Talk to instructors/TAs if you have any questions

[1] <https://engineering.buffalo.edu/computer-science-engineering/information-for-students/policies/academic-integrity.html>

[2] <https://academicintegrity.buffalo.edu/policies.php>

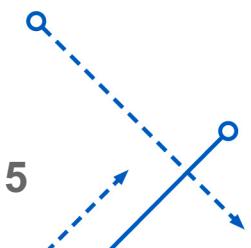


# Questions/Concerns?

# CSE 220: Expectations

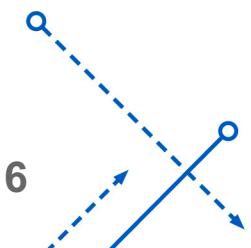
- Systems hacking can be fun!
- My best undergrad course was Operating Systems – the sort-of follow up to this course
- You get out of it what you put in; no more, no less !
- If you are willing to put in time, we are more than happy to help
- If you need a better grade, do better work

I hate grade negotiations at the end of the semester – please chat with me through the semester on your progress, but you will get the grade you earned at the end !



# CSE 220: Other Logistics

- Re-grading is done only for grading errors
- No incompletes are given in the class
- In principle, no makeup exams will be given except for valid reasons  
Please make sure you talk to me well in advance if you have a valid reason – I will not entertain last minute requests
- **No grades will be changed** for any reason other than grading error



# CSE 220: Course Materials

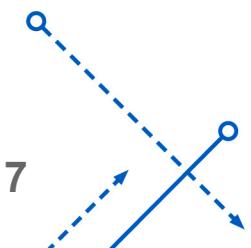
- Follow progress on the course website
- Revise using slides on the course website
- Each class has required and optional readings – read them even if I don't mention it in class
- All discussion/communication will be via Piazza

Register on Piazza today!

Make sure you read the posts before you ask a question

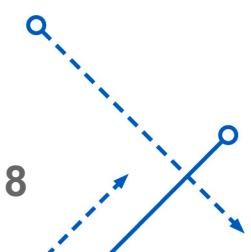
However, if you are not sure about something, just ask!

All important class announcements and materials will be posted on Piazza – you are responsible to follow them!



- You will work on Linux on x86-64 hardware
- We have a VM Image[1] for you to get started quickly
- You don't need to use this image
  - But you will be responsible to make sure your submissions are compatible to this image
  - We will not support platform issues on any other platform
- If in doubt, please use this image

[1] <https://www.cse.buffalo.edu/~eblanton/misc/vm/>

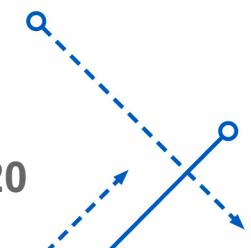


# Programming Projects

- Significant portion of your course grade will be projects  
These are individual projects  
Projects will be written in C
- To reiterate: projects must run on the course VM image
- We will use GitHub Classroom  
For assignment distribution  
For providing assistance
- You must have (or create) a GitHub account
- You are expected to use git and GitHub for development
- If you are not already familiar, learn git
  - [1] Git book: <https://git-scm.com/book/en/v2>
  - [2] Git tutorial: <https://alistapart.com/article/get-started-with-git/>
  - [3] Git usage: <https://www.google.com/search?q=using+git>

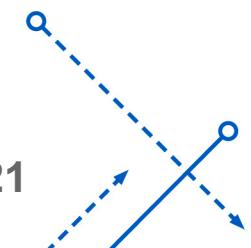
# Project Assistance

- TAs will be your primary source of help for projects
- To get the most out of the TAs, do:
  - Try the obvious things first
  - Create minimal examples to show problems
- Consult the documentation
- To avoid wasting TA time and failing to get help, don't:
  - Ask for help before you've tried to understand the problem
  - Start at the last minute
- That said, if you've spent many hours to identify the problem and can't, please ask for help



# Programming Tools

- We will learn a lot of tools in this class
- You will be expected to use a few tools for this course:  
The C compiler  
Make  
Gdb  
A programming editor  
Others ...
- We will help you learn these tools



# Editors

- We don't care what editor you use
- Your life will be simpler if you use a programmer's editor that can do the following:
  - Syntax highlighting
  - Automatic indentation
  - Brace/paranthesis matching
  - Extensibility

- It will be difficult for us to help you if you are not using a reasonable environment

- I personally use emacs – installed on the course VM image

- Other candidates

vi

Sublime

Atom

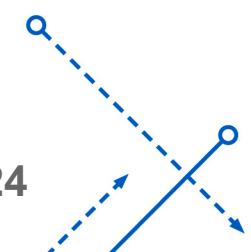
# Project Submission

- All submissions are done through autograder [1]
- Submission rules:
  - Submitted within 24 hrs after the deadline: -20%
  - Doesn't count Saturday or Sunday
  - Doesn't count University holidays
- Projects submitted after 24 hrs will not be accepted
- Example: Project is due Friday at 11:59 PM, turned in Monday at 3 pm – 20% penalty
- Example: Project is due Monday at 11:59 PM, turned in Wednesday at 12:15 AM – not accepted

[1] <https://autograder.cse.buffalo.edu/>

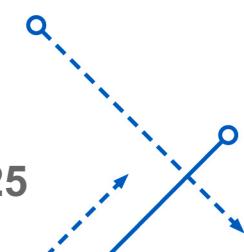
# Grading

- Passing this course requires three major conditions:  
Completion of the AI quiz with perfect score  
Completion of Lab 01 with perfect score  
At least a 60% average on all exams and lab exams
- Failure to achieve any of these three points will result in failure in the course
- Your course grade will be calculated per the information in the syllabus



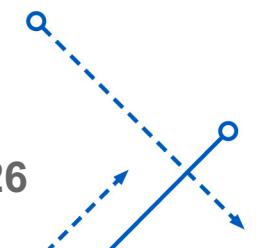
# Today's Assignments

- Immediately:  
Read the Syllabus  
Join the Piazza instance
- By beginning of lab this week:  
Create a GitHub account if you don't already have one  
Download and install the course VM
- By Friday:  
Complete the AI quiz:  
[https://www.cse.buffalo.edu/~eblanton/misc/academic\\_integrity/](https://www.cse.buffalo.edu/~eblanton/misc/academic_integrity/) and  
turn it in on Autograder



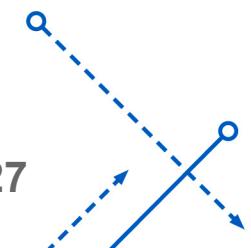
# Labs

We have labs this week!



# Next Class

- Overview of C
- Overview of POSIX API
- Little bit about data types



# Required Readings

- Course syllabus
- K&R: 1.1-1.3

