

Process Layout

Karthik Dantu

Ethan Blanton

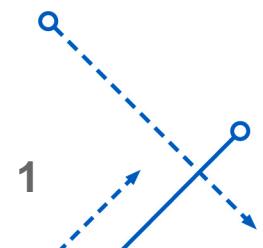
Computer Science and Engineering

University at Buffalo

kdantu@buffalo.edu

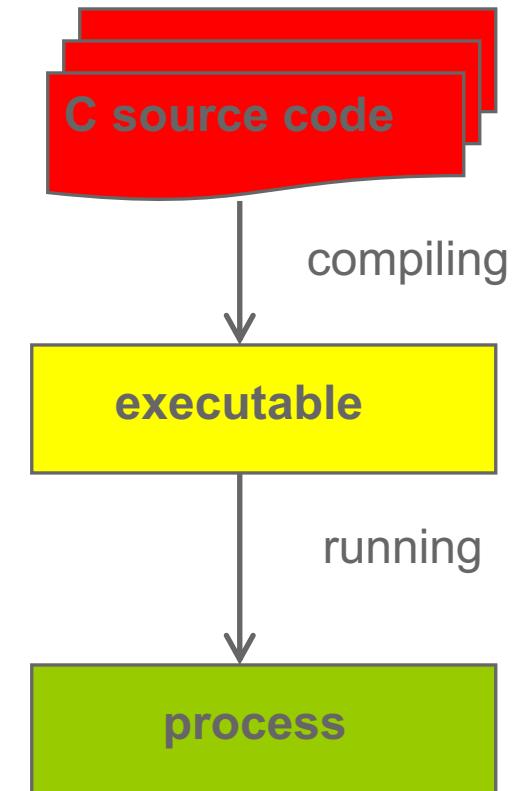
Portions of this lecture are from the Princeton COS 217 course slides

Karthik Dantu



Code → Executable → Process

- **C source code**
C statements organized into functions
Stored as a collection of files (.c and .h)
- **Executable module**
Binary image generated by compiler
Stored as a file (e.g., *a.out*)
- **Process**
Instance of a program that is executing
 - With its own address space in memory
 - With its own id and execution stateManaged by the operating system



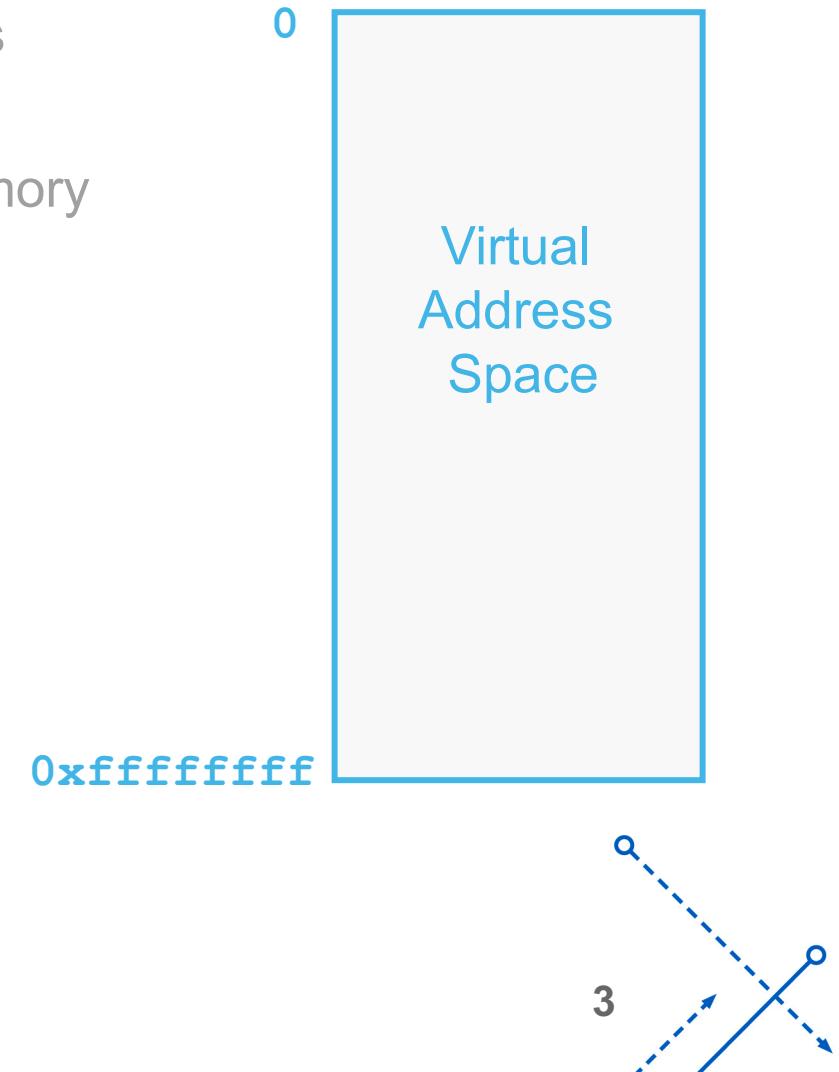
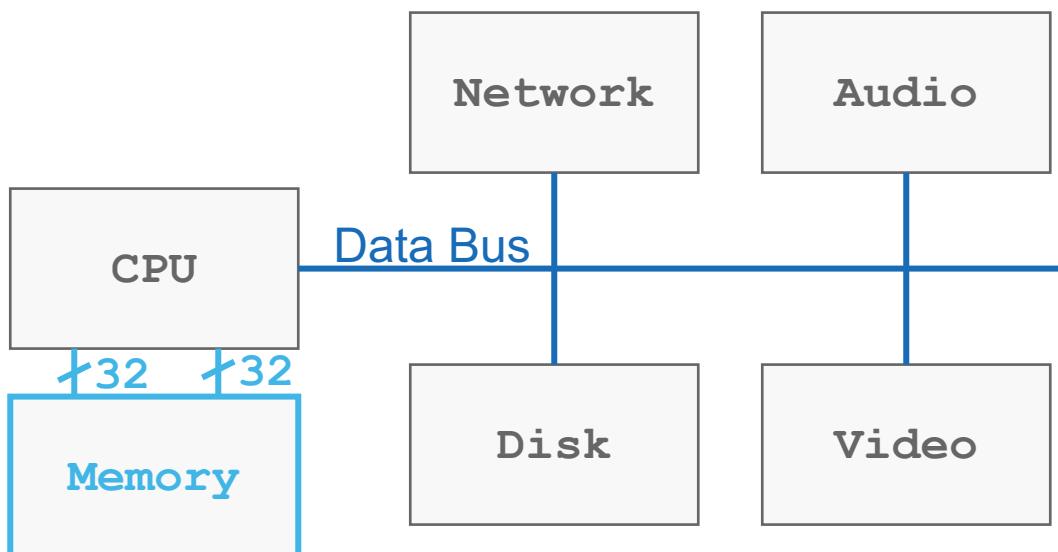
Process Execution

- What is virtual memory?

Contiguous addressable memory space for a single process

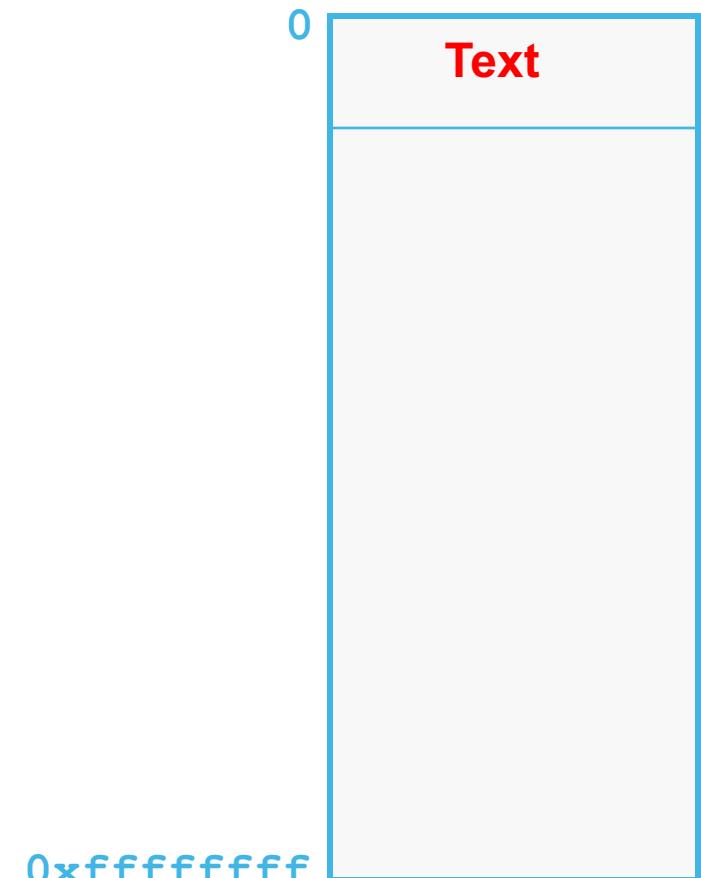
May be swapped into physical memory from disk in pages

Let's you pretend each process has its own contiguous memory



What to Store: Code and Constants

- Executable code and constant data
 - Program binary, and any shared libraries it loads
 - Necessary for OS to read the commands
- OS knows everything in advance
 - Knows amount of space needed
 - Knows the contents of the memory
- Known as the “text” segment
- Note: Some systems (e.g., hats) store some constants in “rodata” section



What to Store: “Static” Data

- Variables that exist for the entire program

Global variables, and “static” local variables

Amount of space required is known in advance

- Data: initialized in the code

Initial value specified by the programmer

E.g., “`int x = 97;`”

Memory is initialized with this value

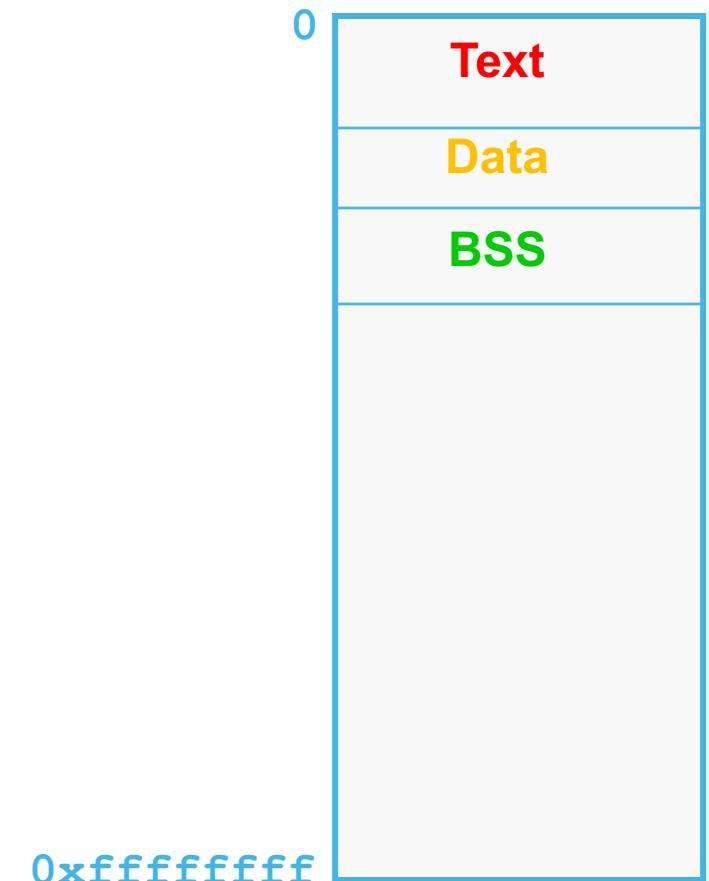
- BSS: not initialized in the code

Initial value not specified

E.g., “`int x;`”

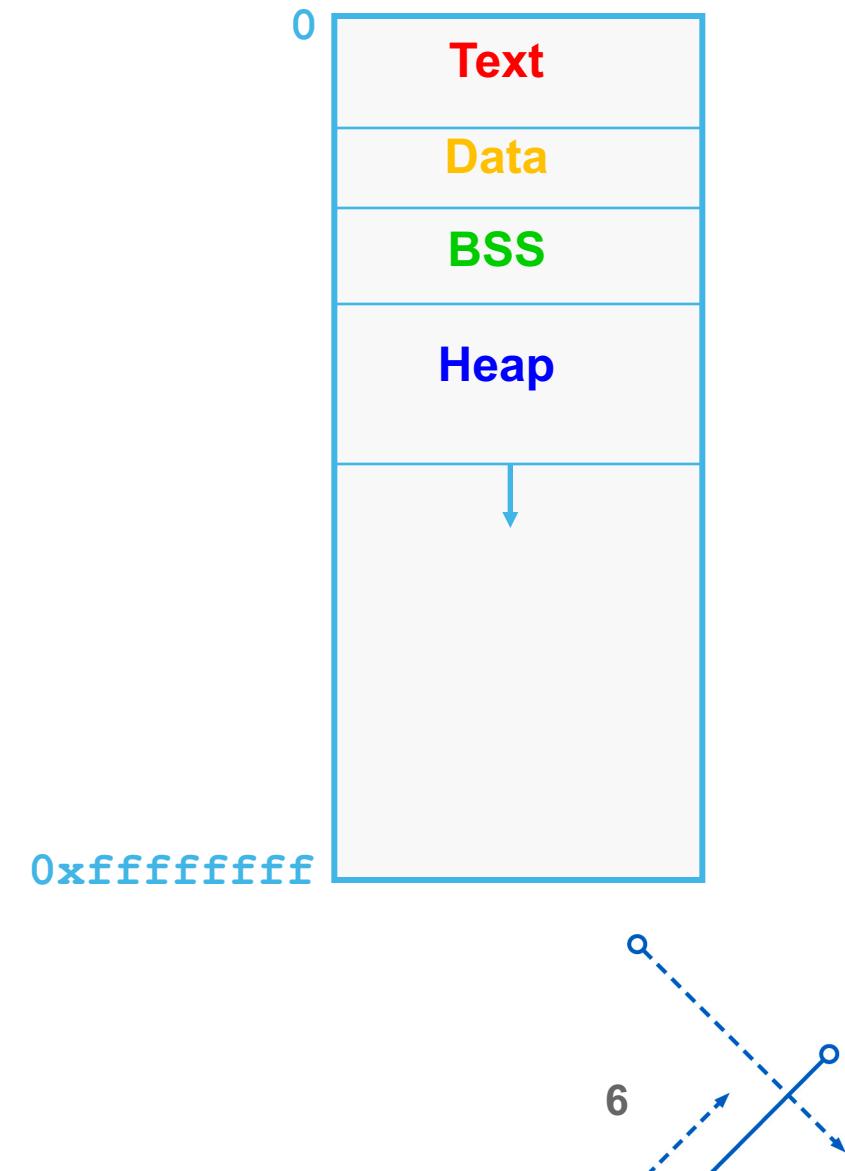
All memory initialized to 0 (on most OS’s)

BSS stands for “Block Started by Symbol”



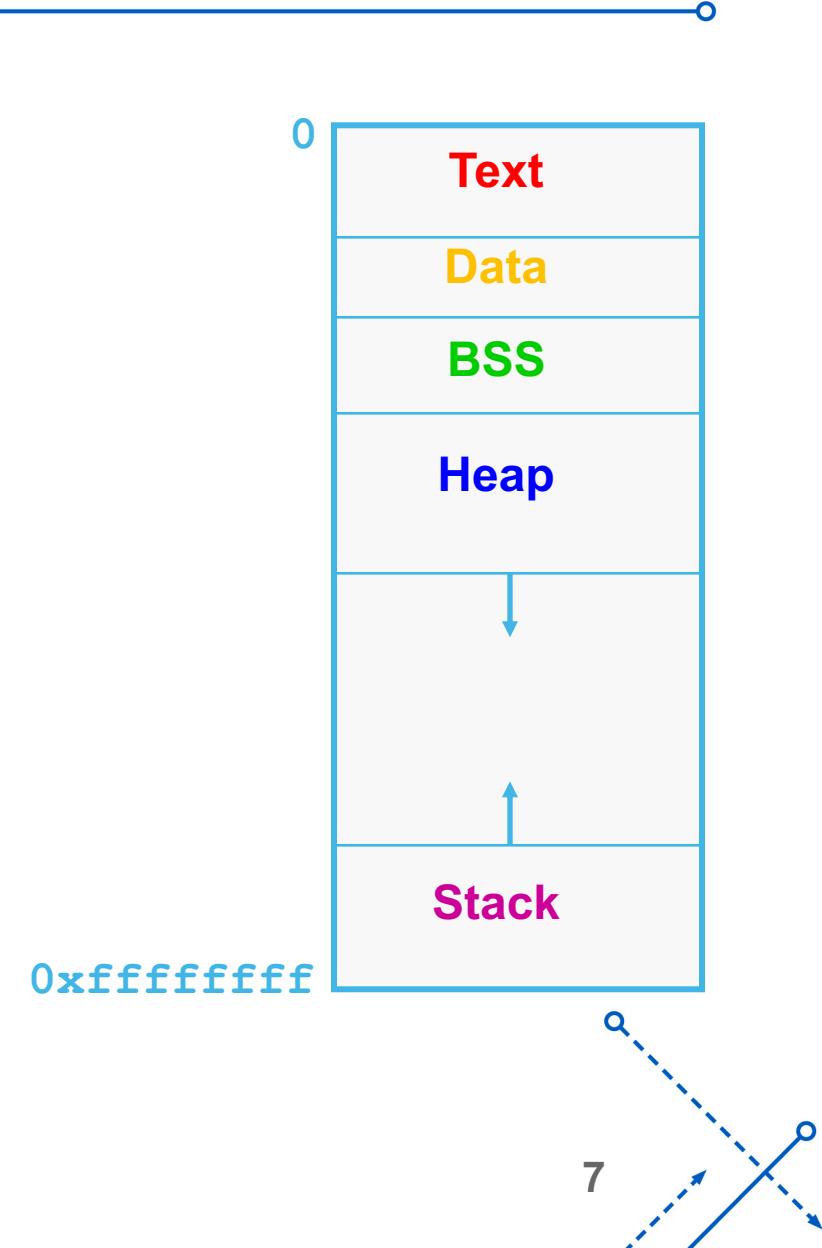
What to Store: Dynamic Memory

- Memory allocated while program is running
 - E.g., allocated using the `malloc()` function
 - And deallocated using the `free()` function
- OS knows nothing in advance
 - Doesn't know the amount of space
 - Doesn't know the contents
- So, need to allow room to grow
 - Known as the “**heap**”
 - Detailed example in a few slides
 - More in programming assignment #4



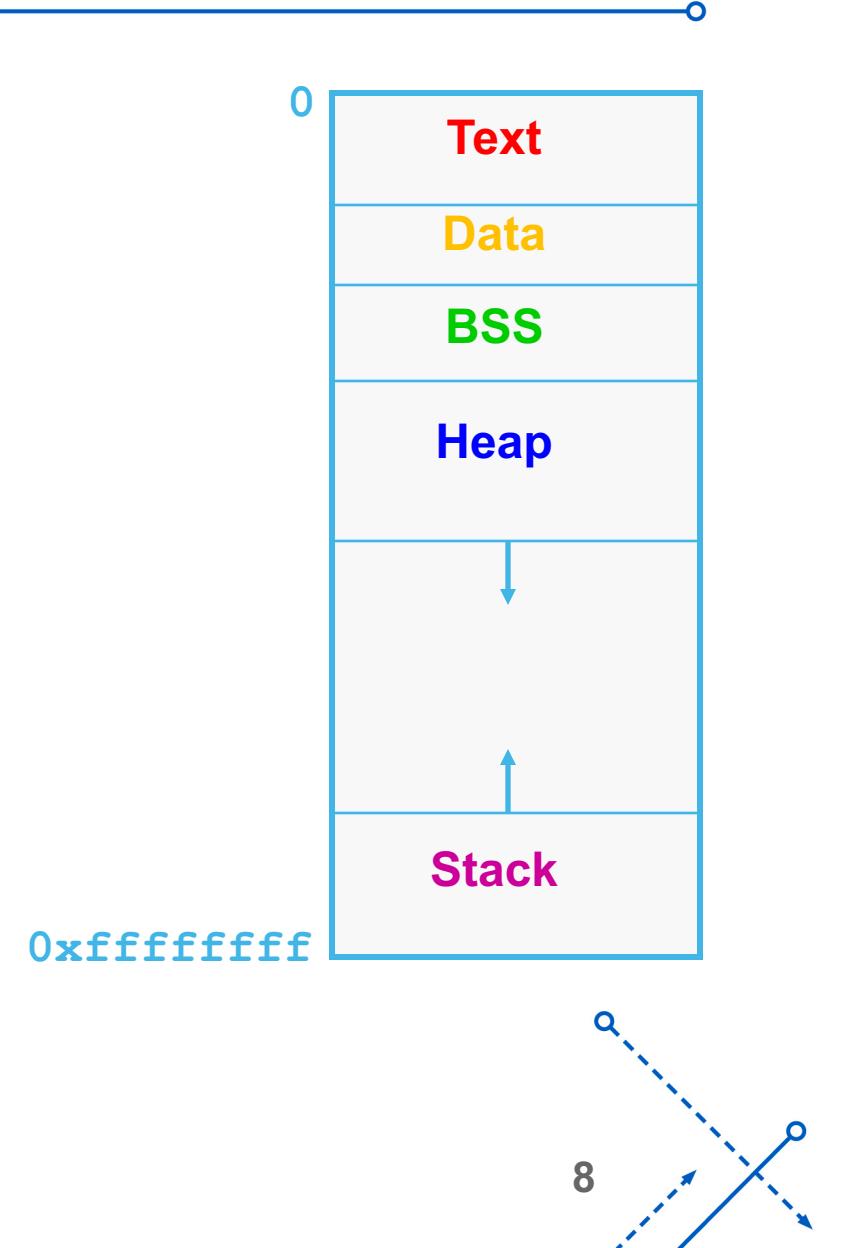
What to Store: Temporary Variables

- Temporary memory during lifetime of a function or block
 - Storage for function parameters and local variables
- Need to support nested function calls
 - One function calls another, and so on
 - Store the variables of calling function
 - Know where to return when done
- So, must allow room to grow
 - Known as the “stack”
 - Push on the stack as new function is called
 - Pop off the stack as the function ends
- Detailed example later on



Memory Layout: Summary

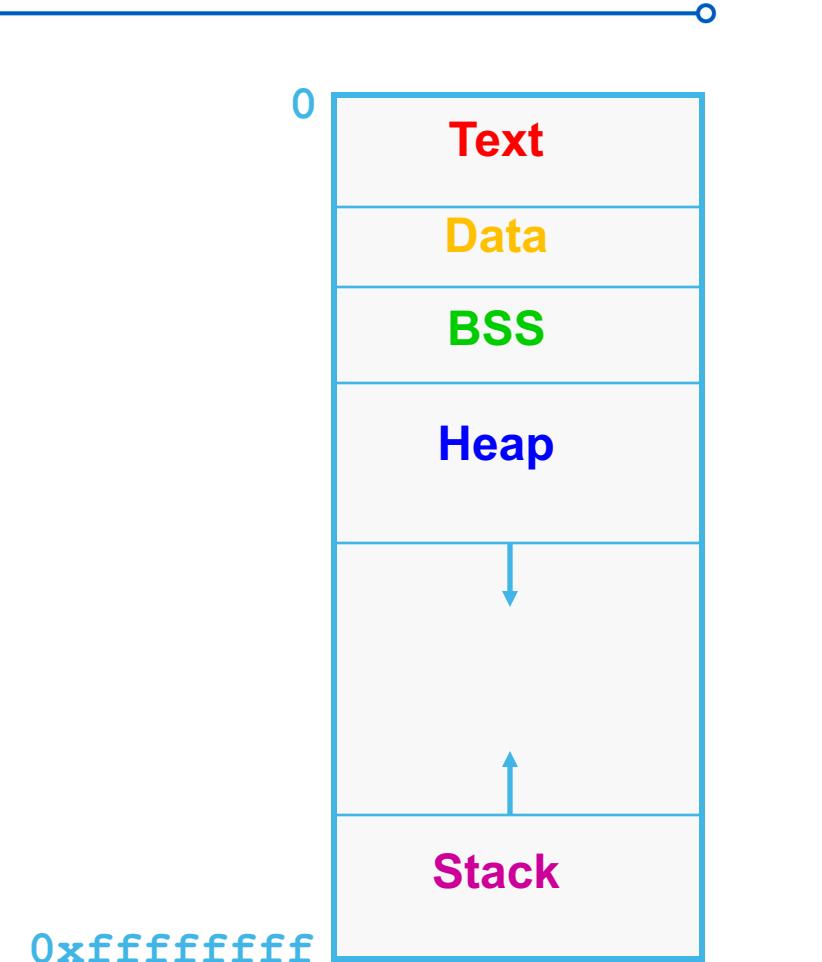
- **Text**: code, constant data
- **Data**: initialized global & static variables
- **BSS**: uninitialized global & static variables
- **Heap**: dynamic memory
- **Stack**: local variables



Memory Layout: Example

```
char* string = "hello";
int iSize;

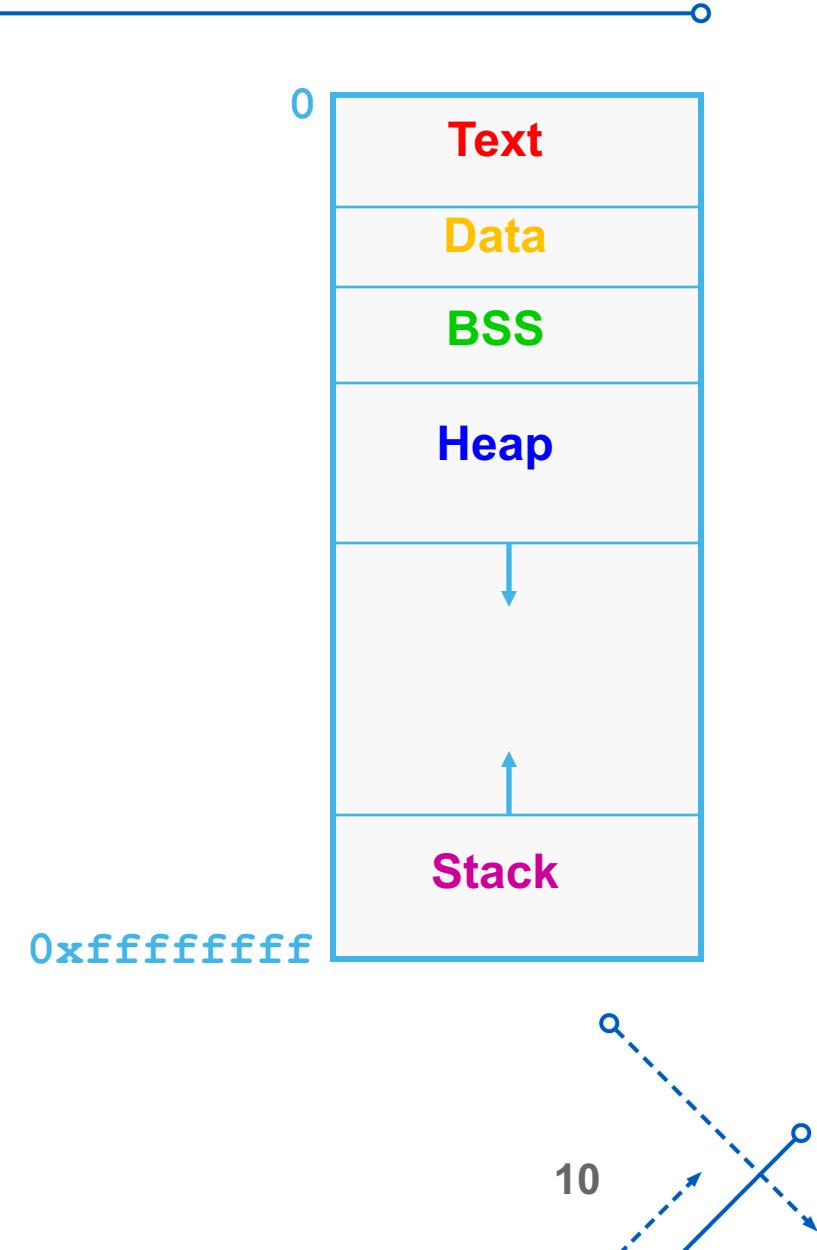
char* f(void)
{
    char* p;
    iSize = 8;
    p = malloc(iSize);
    return p;
}
```



Memory Layout: Example

```
char* string = "hello";
int iSize;

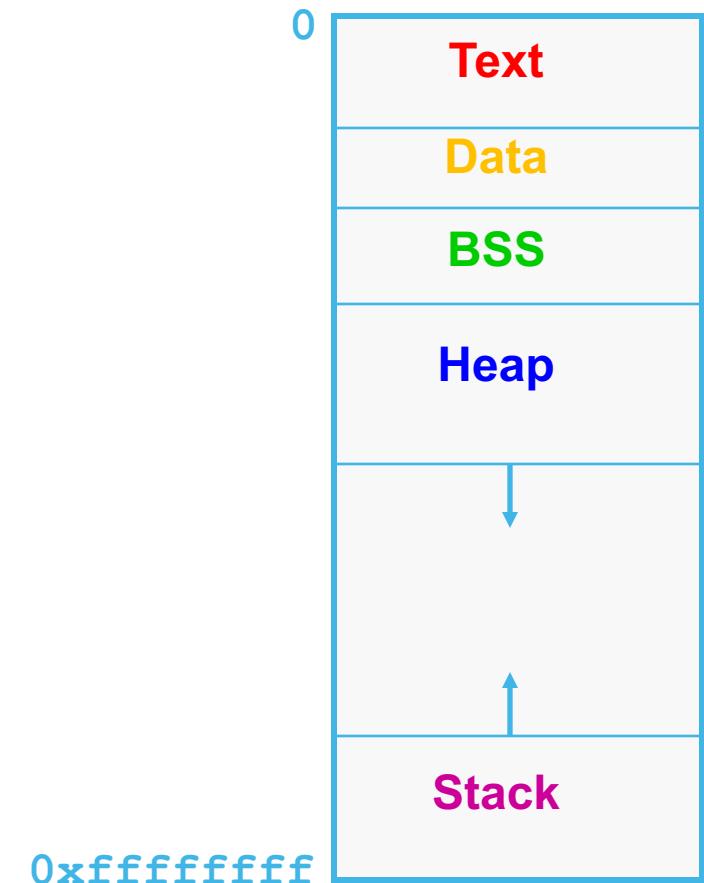
char* f(void)
{
    char* p;
    iSize = 8;
    p = malloc(iSize);
    return p;
}
```



Memory Layout: Data

```
char* string = "hello";
int iSize;

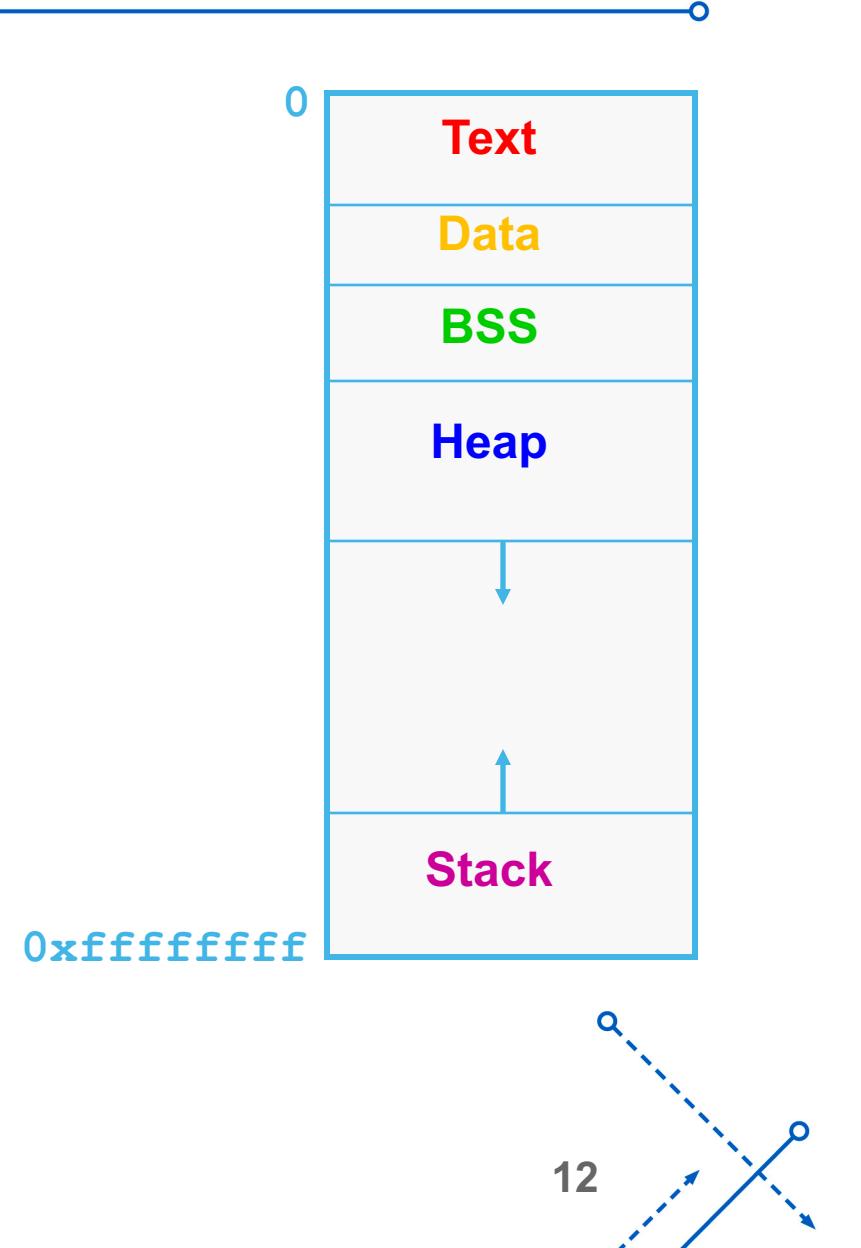
char* f(void)
{
    char* p;
    iSize = 8;
    p = malloc(iSize);
    return p;
}
```



Memory Layout: BSS

```
char* string = "hello";
int iSize;

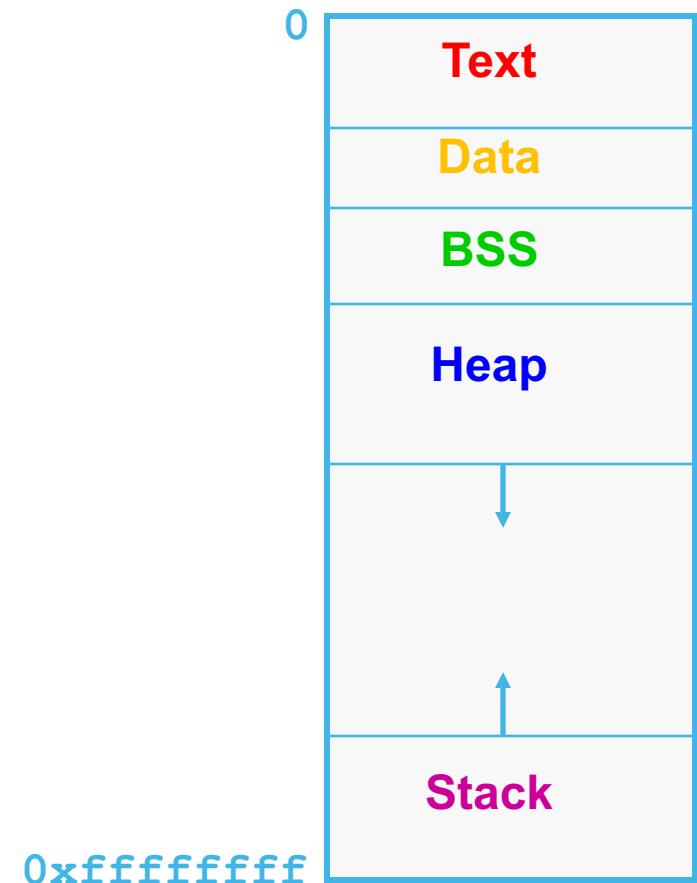
char* f(void)
{
    char* p;
    iSize = 8;
    p = malloc(iSize);
    return p;
}
```



Memory Layout: Heap

```
char* string = "hello";
int iSize;

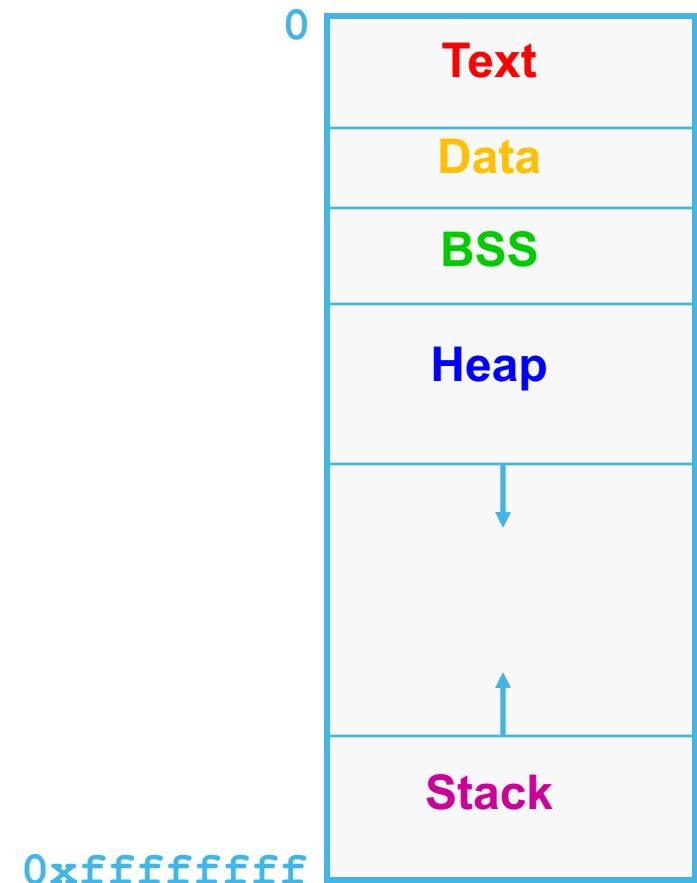
char* f(void)
{
    char* p;
    iSize = 8;
    p = malloc(iSize);
    return p;
}
```



Memory Layout: Stack

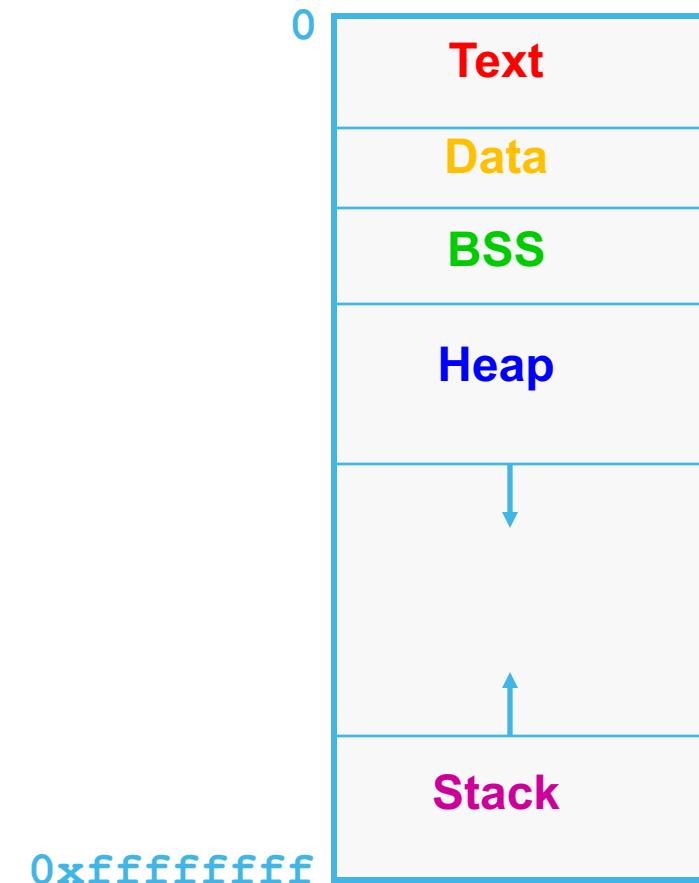
```
char* string = "hello";
int iSize;

char* f(void)
{
    char* p;
    iSize = 8;
    p = malloc(iSize);
    return p;
}
```



Memory Allocation and De-allocation

- How, and when, is memory allocated?
 - Global and static variables: program startup
 - Local variables: function call
 - Dynamic memory: `malloc()`
- How is memory deallocated?
 - Global and static variables: program finish
 - Local variables: function return
 - Dynamic memory: `free()`
- All memory deallocated when program ends
 - It is good style to free allocated memory anyway



Memory Allocation Example

```
char* string = "hello"; ← Data: "hello" at startup
int iSize; ← BSS: 0 at startup
```

```
char* f(void)
{
    char* p;
    iSize = 8;
    p = malloc(iSize);
    return p;
}
```

Stack: at function call

Heap: 8 bytes at malloc

Memory Deallocation Example

```
char* string = "hello";
```

Available till termination

```
int iSize;
```

Available till termination

```
char* f(void)
```

```
{
```

```
    char* p;
```

```
    iSize = 8;
```

```
    p = malloc(iSize);
```

```
    return p;
```

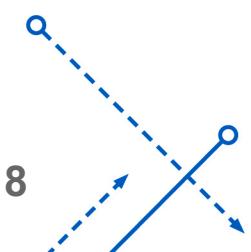
Deallocate on return from f

Deallocate on free()

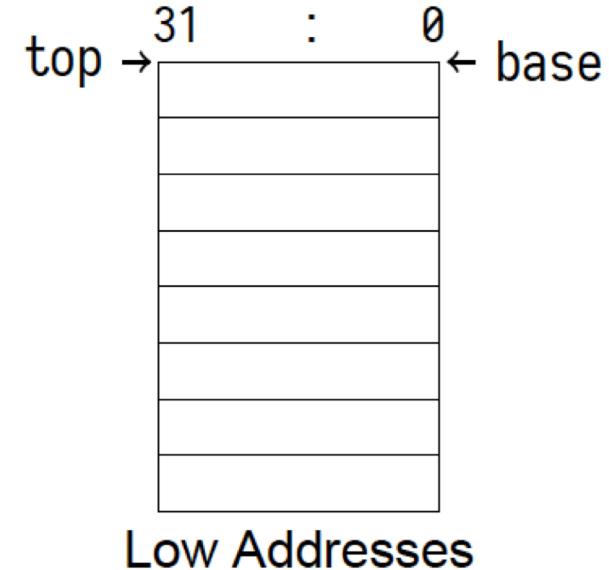
```
}
```

Aside: Using Sections

- The exact addresses of sections will vary
- However, you can usually assume certain things
- We'll look at some of those properties later
- Learning to recognize the location of a pointer is valuable
- For example: all pointers < 4096 (0x1000) are invalid!

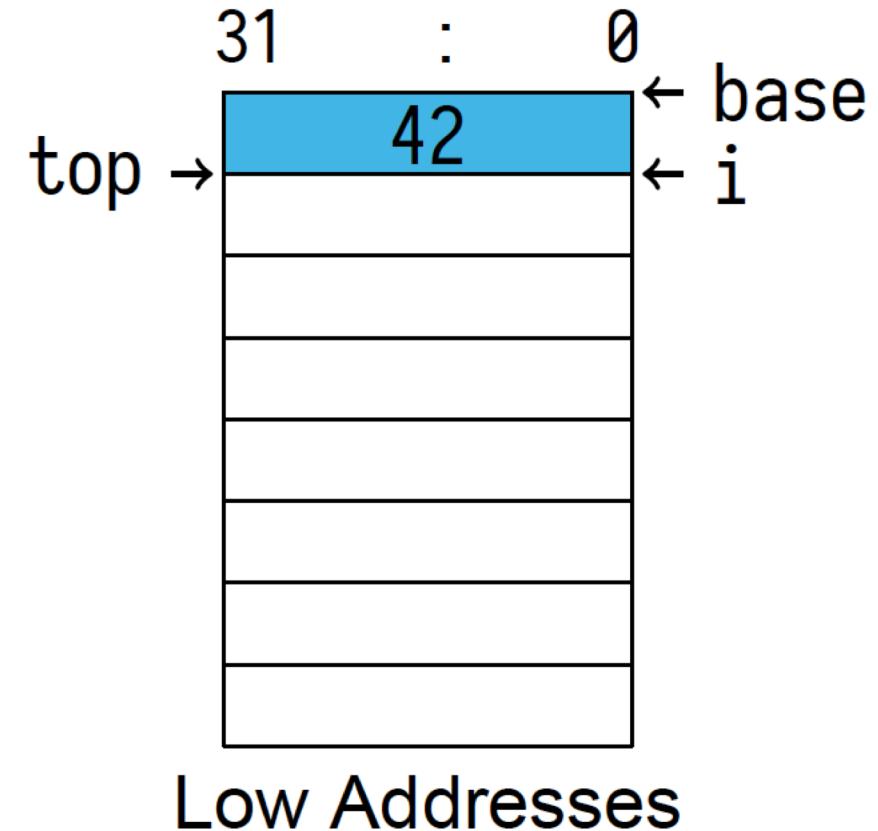


Stack Operations

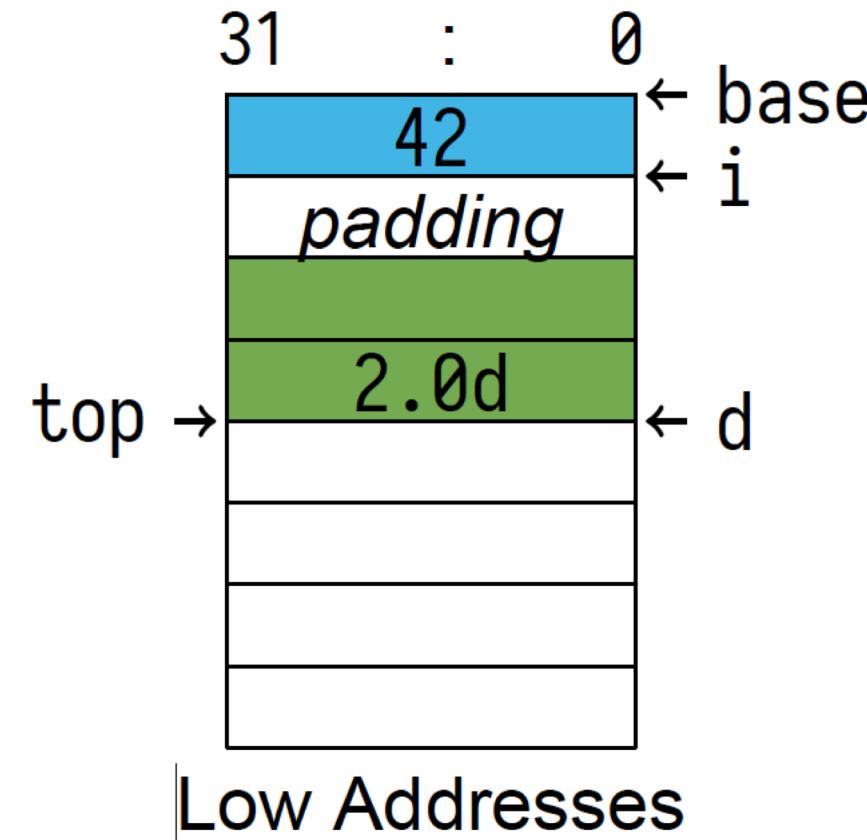


(An empty stack; each row is 32 bits.)

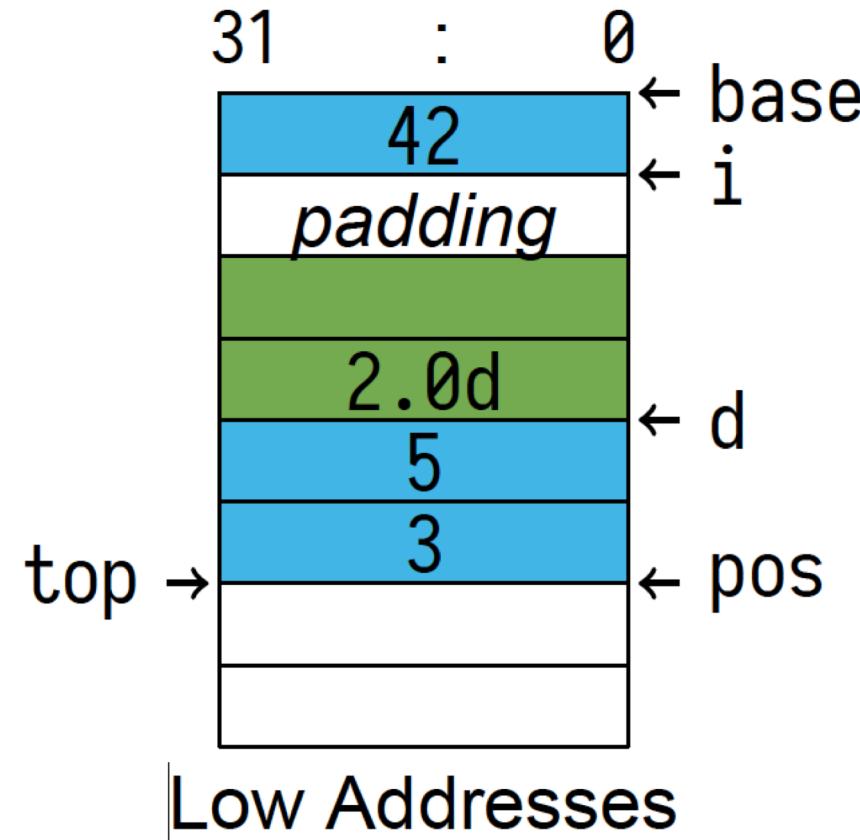
Stack Operations



```
push int i = 42;
```

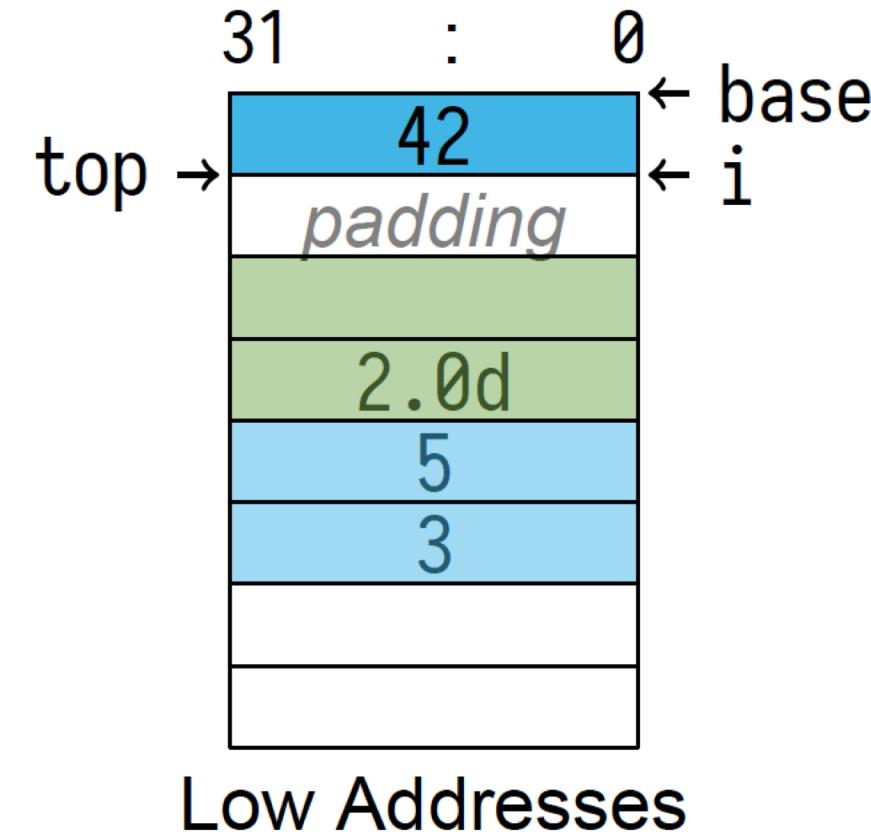


```
push double d = 2.0;  
(Remember padding!)
```



```
push struct { int x; int y; } pos = { x = 3, y = 5 };
```

Stack items are typically referenced with respect to its *top*.
E.g., *d* is at *top* + 8



pop 20 bytes to remove pos and d

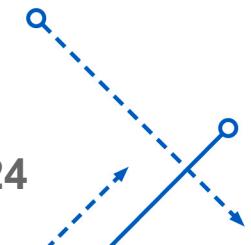
Note that the unused data remains present on the stack.

Variable Declarations

- A variable does two things
 - Ask compiler to reserve memory for data
 - Name the location of that data

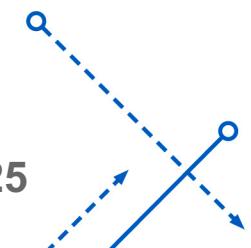
```
int array[32];
```

- “Make space for 32 integers and call that space array
- Every non-static, local variable is an automatic variable



Automatic Variable Lifetime

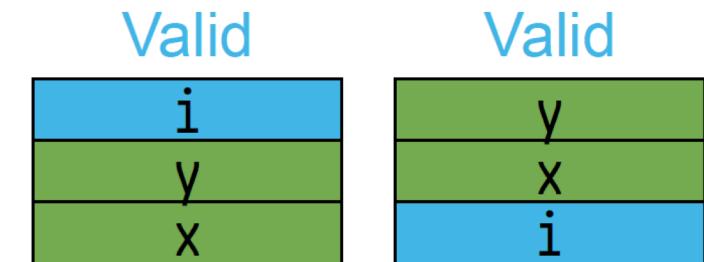
- Automatic variables are:
 - Guaranteed to be allocated before they are first referenced
 - Guaranteed to be valid until their enclosing block is done
- In many cases they are created when the function is entered
- Placing automatic variables on the stack allows this



Automatic Variable Placement

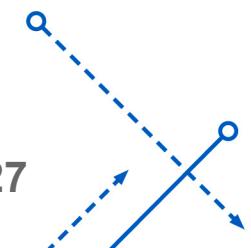
- Automatic variables may be allocated anywhere
- The programmer cannot predict their order or location
- They may only be in registers!
- Their structure will be preserved

```
int i;  
struct {  
    int x; int y;  
} pos;
```



Function Call Nesting

- Note that:
 - Function calls form a tree over the life of a program
 - Function calls form a stack at any point in time
- This is because:
 - A function may call many functions consecutively
 - A function can call only one function at a time
- These properties directly affect the program stack



Function Calls

- At its simplest, a function call consists of:
 - A jump to a new program location
 - Execution of the function code
 - A jump back to the calling location
- However, many function calls are more complicated. They may:
 - Allocate automatic variables
 - Call other functions
 - Temporarily save registers
 - ...
- In these cases, functions require a stack frame.



Stack Frames

- A stack frame holds information for a single function invocation.
- While the details vary by platform, it will include:
 - Saved processor registers
 - Local variables for the current function
 - Arguments for any called function
 - The return location for any called function
- We will discuss all of these except saved processor registers.

(Maybe we'll get to those later.)



Local Variables

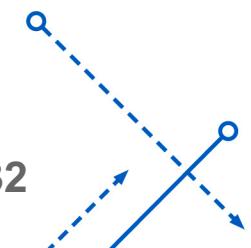
- We have previously discussed automatic variables.
- Often, all local variables for a function are allocated together.
- When the function is entered, it will immediately move the top of the stack to make room for its local storage.
- This portion of the stack frame is then of fixed size.
- Its size is often not saved, but recorded in the program instructions by the compiler.
- The location of individual variables are likewise recorded.

Function Arguments

- The platform ABI will determine how arguments are passed
- Normally, it is a combination of registers and stack space
- On x86-64 Linux, the first six 64 bit values are passed in registers
- Any additional arguments are pushed onto the stack
- Therefore, many functions have no arguments on the stack

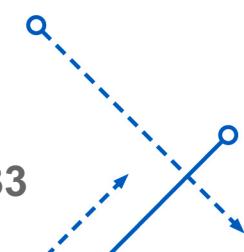
Function Arguments Layout

- If function arguments are pushed onto the stack, they are normally pushed in reverse order
- That is, the first function argument is closest to the top
- Among other reasons, this allows for a variable number of arguments
- Consider `printf`: it takes 1 or more arguments
- The first format argument tells it how many

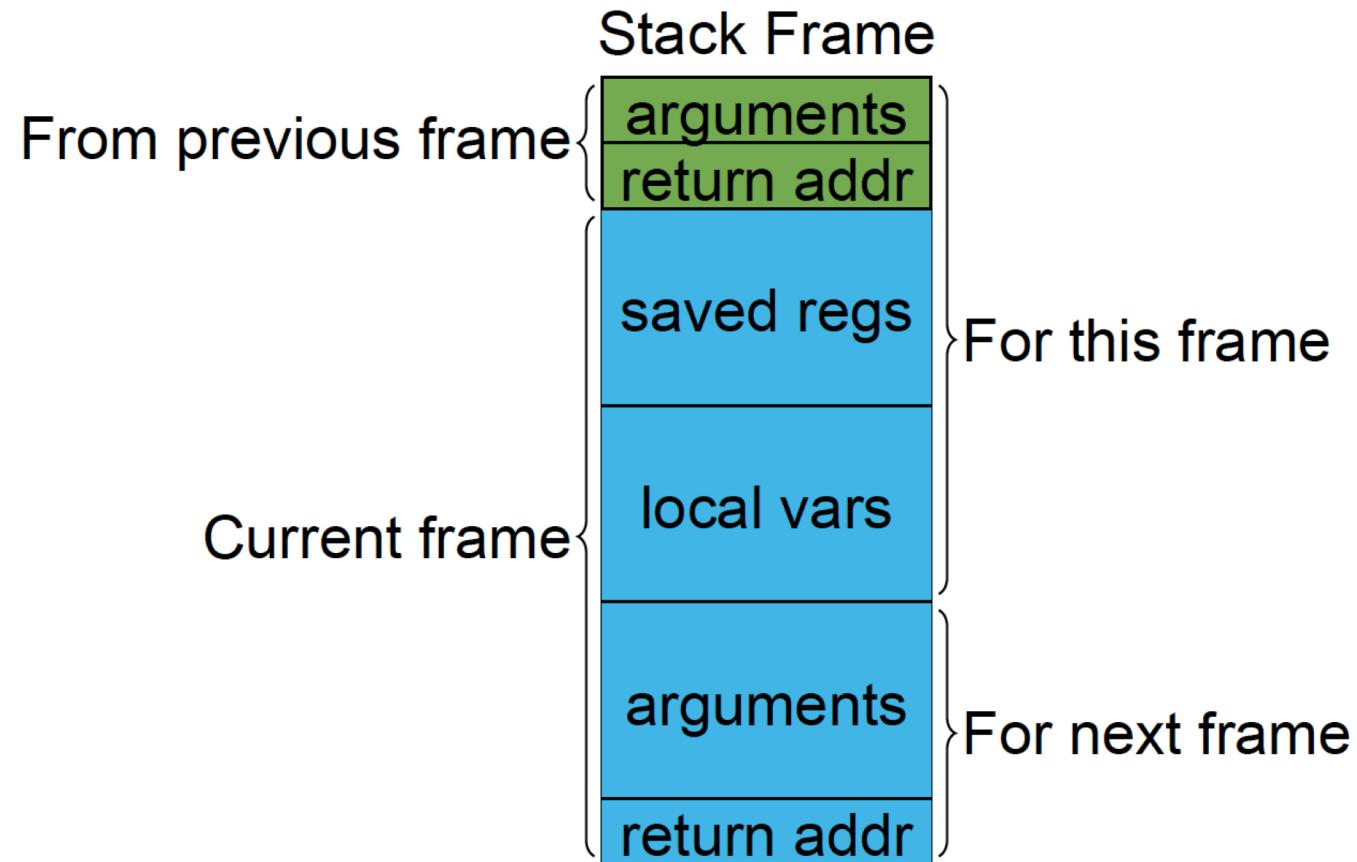


The Program Counter

- The other major item that must be tracked for the function call stack is the program counter
- The program counter is the address of the machine instruction the processor is currently executing
- For a function call:
 - the current program counter is pushed before jumping to the called function
 - the called function pops the program counter in order to return
- On some architectures there is a dedicated instruction for this

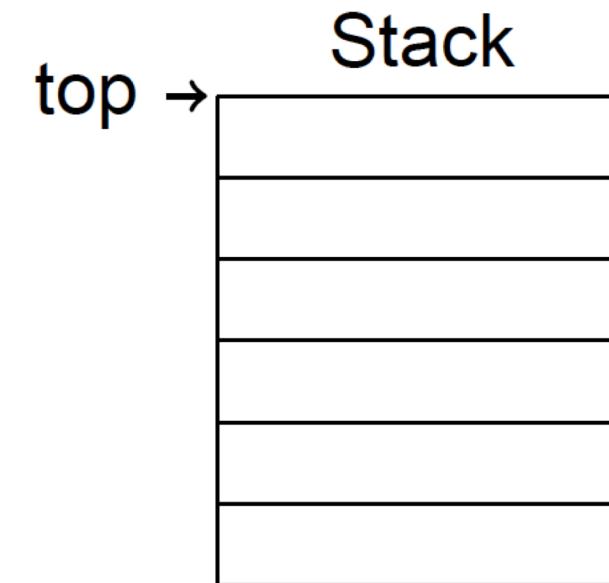


A stack frame



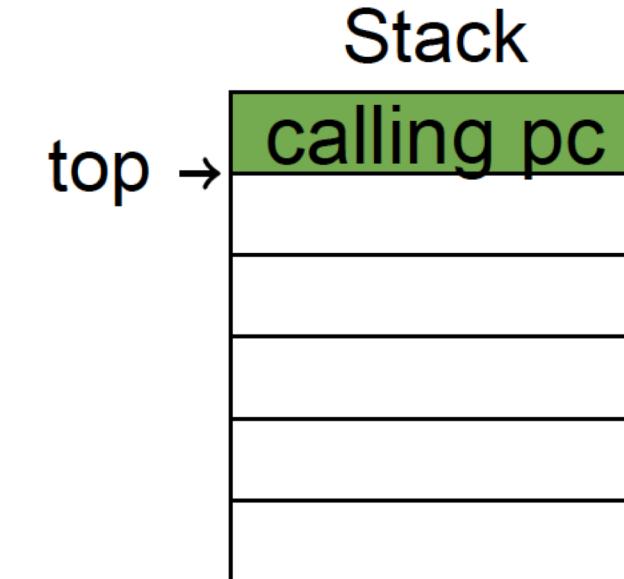
Stack Frame: Example

```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



Stack Frame: Example

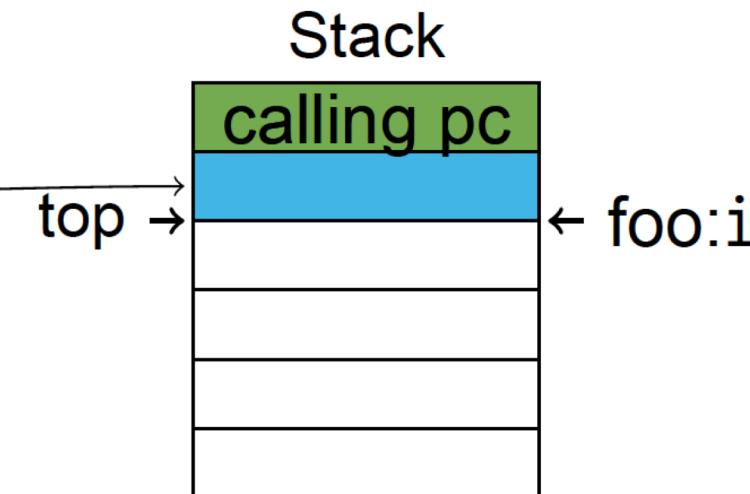
```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



call foo()

Stack Frame: Example

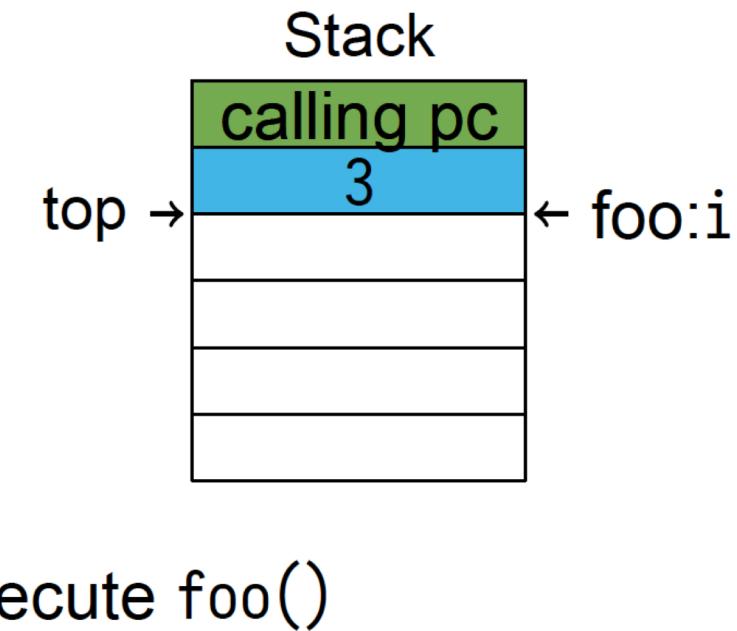
```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



Reserve space for foo()'s locals

Stack Frame: Example

```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



Stack Frame: Example

```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```

Stack

calling pc
3
3

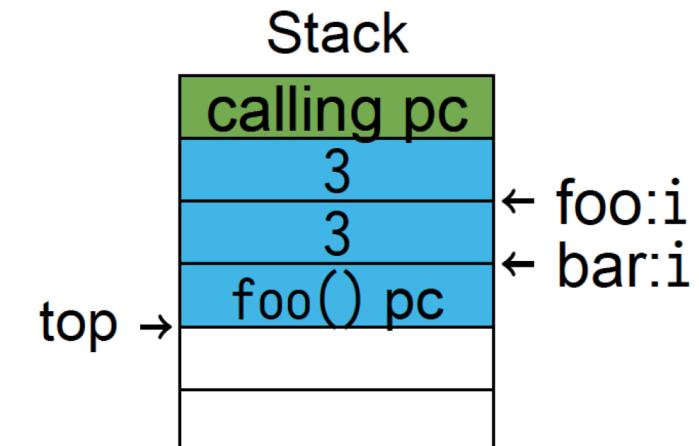
← foo:i
← bar:i

top →

Execute foo();
prepare to call bar()

Stack Frame: Example

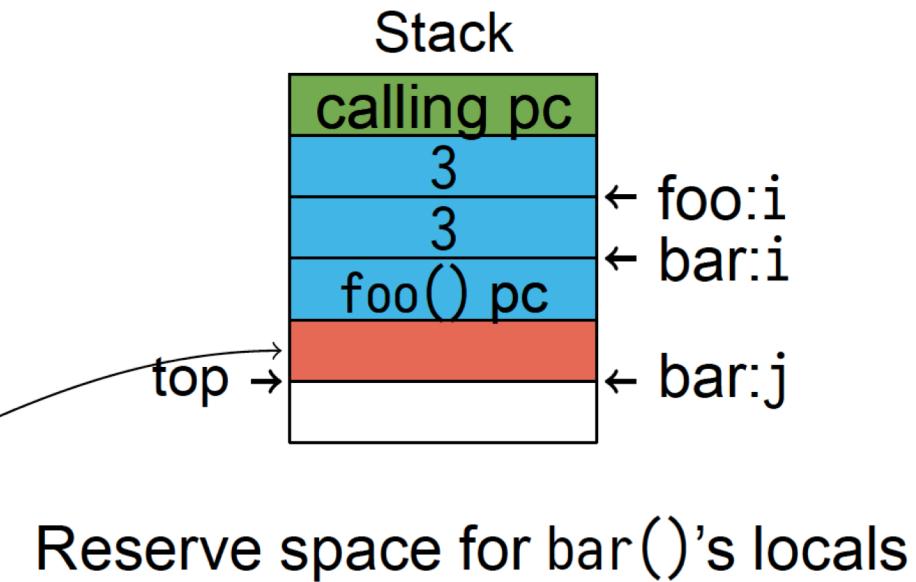
```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



Push PC; call bar()

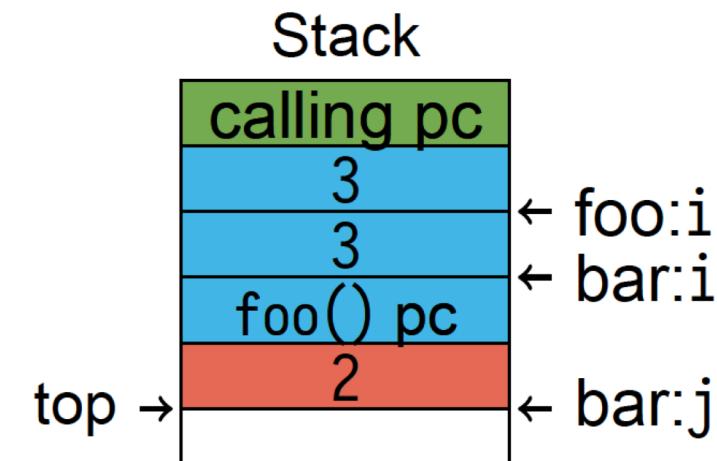
Stack Frame: Example

```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



Stack Frame: Example

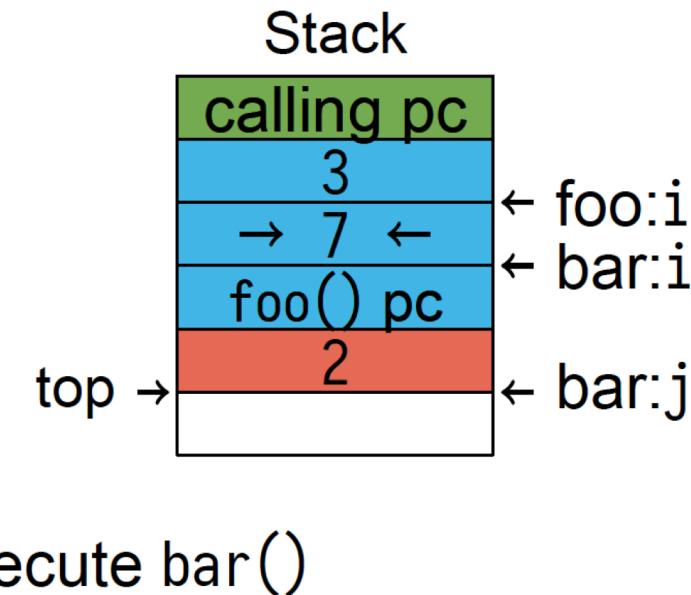
```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
    i = 5 + j;  
}
```



Execute bar()

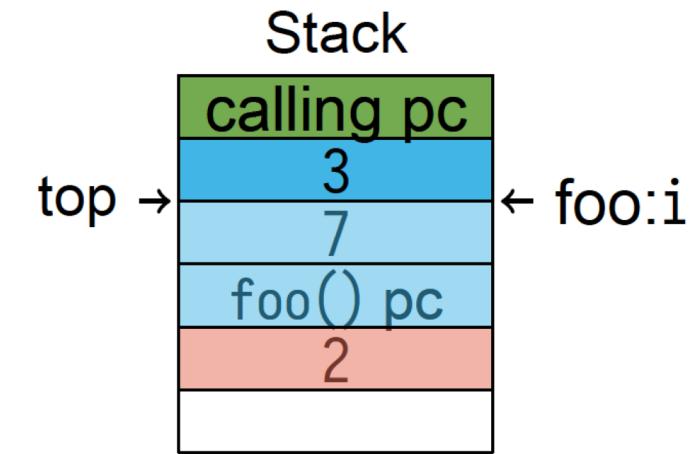
Stack Frame: Example

```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



Stack Frame: Example

```
void foo() {  
    int i = 3;  
  
    bar(i);  
    /* ... */  
}  
  
void bar(int i) {  
    int j = 2;  
  
    i = 5 + j;  
}
```



Return from bar();
Pop bar()'s stack frame;
Execute foo()

Summary

- POSIX programs are laid out in sections
The stack is a section
- The stack grows downward
- Automatic variables are allocated on the stack
- Stack frames track function calls
- Items removed from the stack are not cleared
- Stack-allocated arguments are why C is call-by-value

